

CS2109S: Introduction to AI and Machine Learning

Lecture 9: Intro to Neural Networks

20 Oct 2025

Outline

- Neural Network
 - Neuron
 - Activation Functions
 - Multi-layer Neural Network
- Tasks for Neural Network
 - Binary Classification
 - Multi-class Classification
 - Single-output Regression
 - Multi-output Regression
- Backpropagation

Outline

- **Neural Network**
 - Neuron
 - Activation Functions
 - Multi-layer Neural Network
- Tasks for Neural Network
 - Binary Classification
 - Multi-class Classification
 - Single-output Regression
 - Multi-output Regression
- Backpropagation

Logistic Regression Model

A logistic regression model computes a weighted sum of input features and passes the result through a **sigmoid function**.

$$z = \sum_{j=0}^d w_j x_j$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

Here, $w_0 \dots w_d$ and $x_0 \dots x_d$ are weights and input features respectively. d is the dimension of the input features.

Linear Regression Model

A linear regression model computes a weighted sum of input features and passes the result through an identity function.

$$z = \sum_{j=0}^d w_j x_j$$

$$\hat{y} = f(z) = z$$

Here, $w_0 \dots w_d$ and $x_0 \dots x_d$ are weights and input features respectively. d is the dimension of the input features. f represents the identity function

Neuron

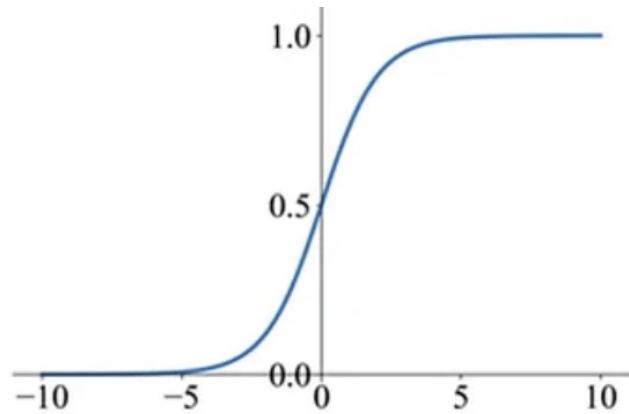
A neuron computes a weighted sum of input features and passes the result through an **activation function**.

$$z = \sum_{j=0}^d w_j x_j$$

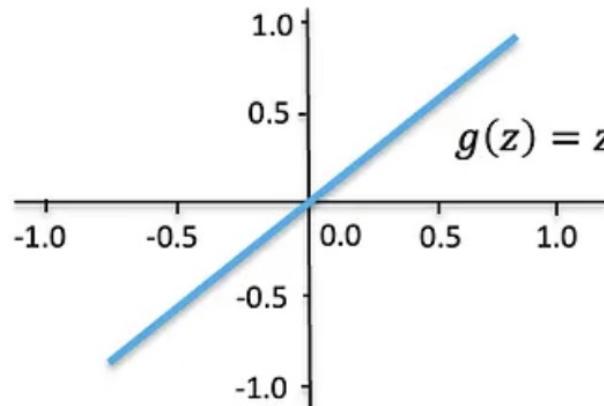
$$\hat{y} = g(z)$$

Here, $w_0 \dots w_d$ and $x_0 \dots x_d$ are weights and input features respectively. d is the dimension of the input features. g represents the activation function

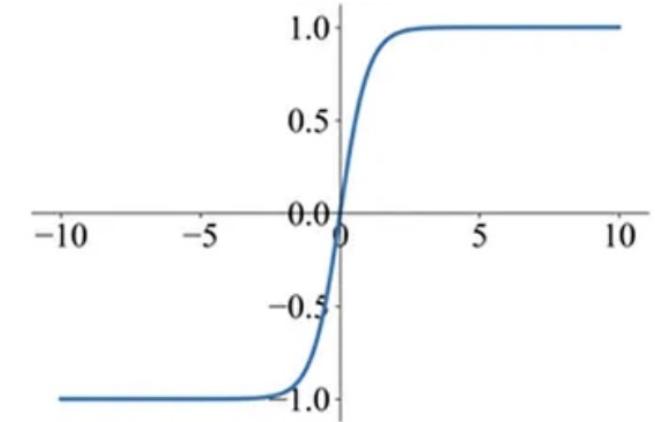
Activation Functions



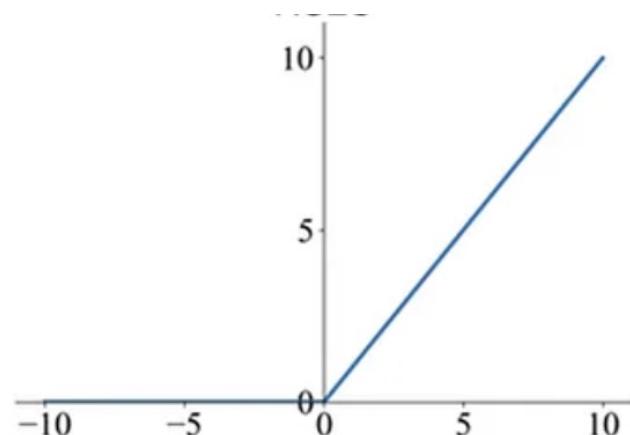
Sigmoid: $g(z) = \sigma(z) = \frac{1}{1+e^{-z}}$



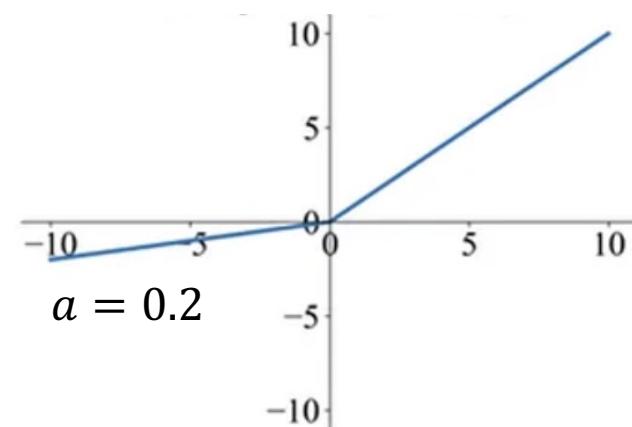
Identity: $g(z) = z$



Tanh: $g(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$



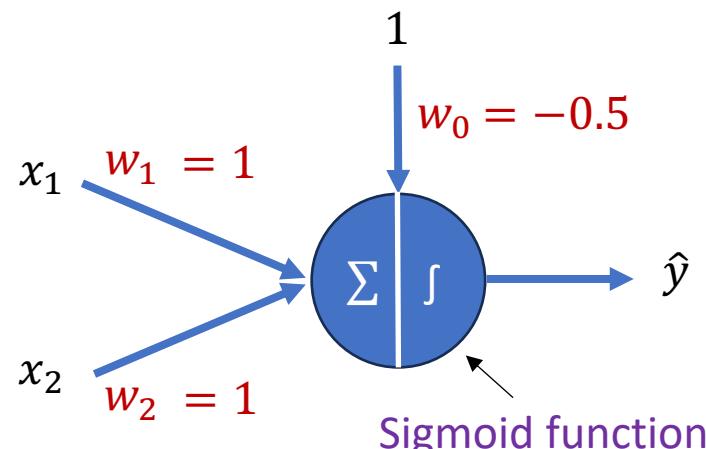
Relu: $g(z) = \max(0, z)$



Leaky Relu: $g(z) = \max(az, z)$

OR Gate Modelling

With One Neuron



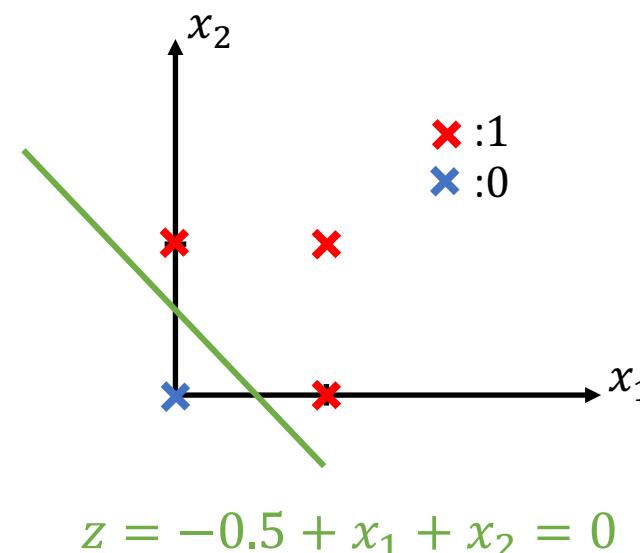
$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z = w_0 + w_1 x_1 + w_2 x_2$$
$$\hat{y} = \sigma(z)$$

Decision threshold: $0.5 = \sigma(0)$

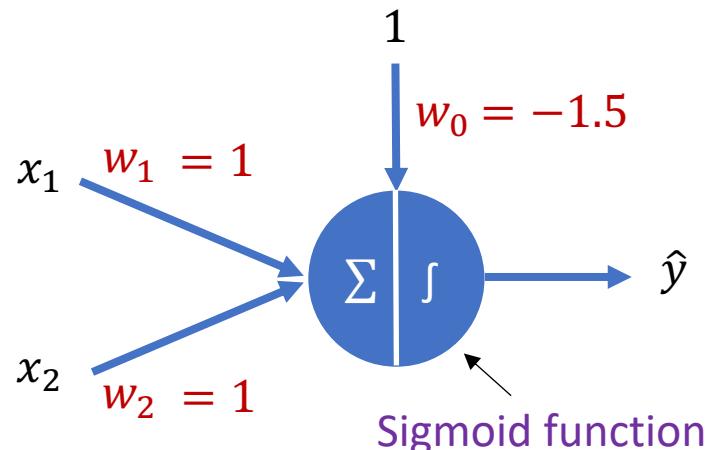
Decision boundary: $z = 0 = w_0 + w_1 x_1 + w_2 x_2$
(A straight line)

OR Gate		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	1



AND Gate Modelling

With One Neuron



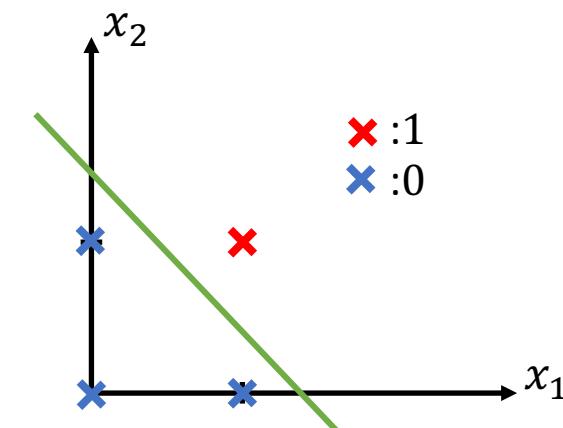
$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z = w_0 + w_1x_1 + w_2x_2$$
$$\hat{y} = \sigma(z)$$

Decision threshold: $0.5 = \sigma(0)$

Decision boundary: $z = 0 = w_0 + w_1x_1 + w_2x_2$
(A straight line)

AND Gate		
x_1	x_2	y
0	0	0
0	1	0
1	0	0
1	1	1

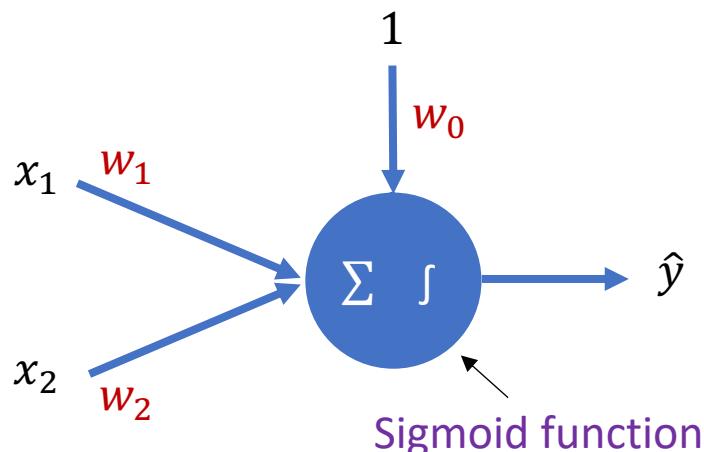


$$z = -1.5 + x_1 + x_2 = 0$$

XNOR Gate Modelling

eXclusive Not OR

With One Neuron



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

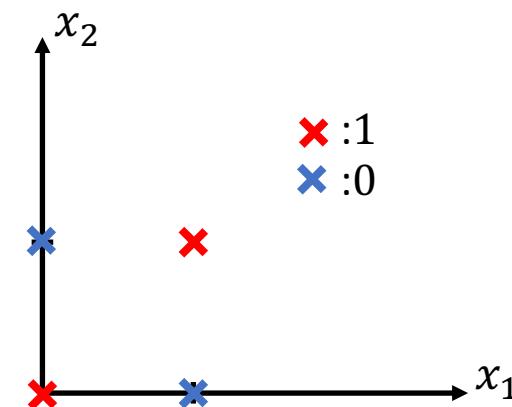
$$z = w_0 + w_1x_1 + w_2x_2$$
$$\hat{y} = \sigma(z)$$

Decision threshold: $0.5 = \sigma(0)$

Decision boundary: $z = 0 = w_0 + w_1x_1 + w_2x_2$
(A straight line)

XNOR Gate

x_1	x_2	y
0	0	1
0	1	0
1	0	0
1	1	1



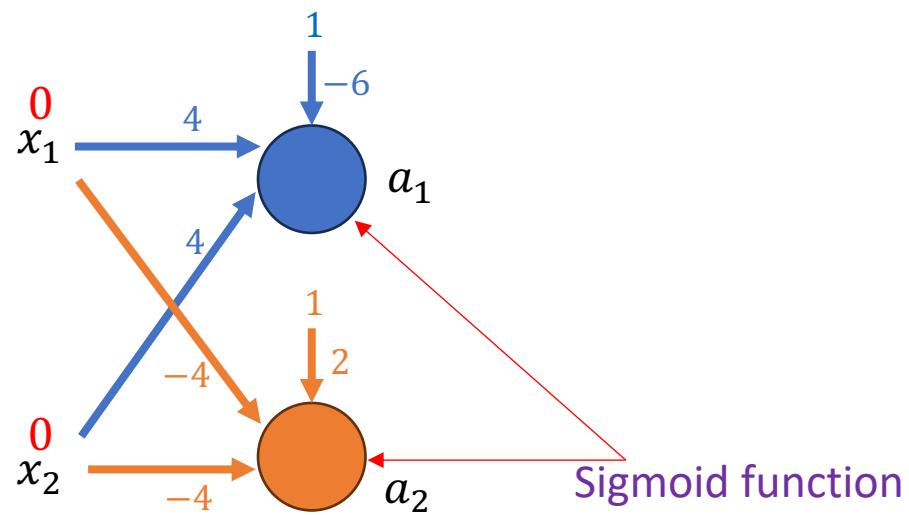
Not linearly separable

Need Feature Transformation

XNOR Gate Modelling

With More Neurons

Neurons can work as learnable feature transformation function.



\$a_1\$ and \$a_2\$ are new features created based on existing features \$x_1\$ and \$x_2\$.

XNOR Gate

\$x_1\$	\$x_2\$	\$y\$	\$a_1\$	\$a_2\$
0	0	1	0.002	0.881
0	1	0		
1	0	0		
1	1	1		

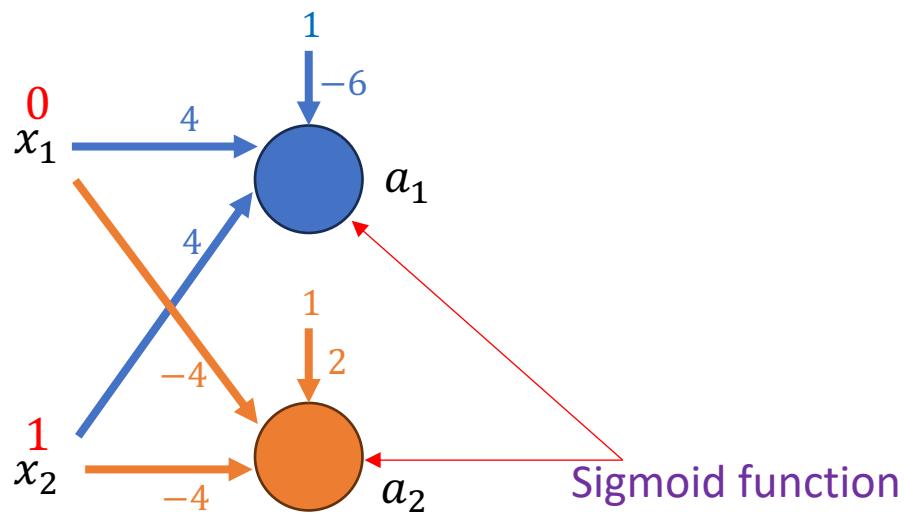
$$\begin{aligned}a_1 &= \sigma(-6 + 4x_1 + 4x_2) \\&= \sigma(-6 + 4 \times 0 + 4 \times 0) = 0.002\end{aligned}$$

$$\begin{aligned}a_2 &= \sigma(2 + (-4)x_1 + (-4)x_2) \\&= \sigma(2 + (-4) \times 0 + (-4) \times 0) = 0.881\end{aligned}$$

XNOR Gate Modelling

With More Neurons

Neurons can work as learnable feature transformation function.



a_1 and a_2 are new features created based on existing features x_1 and x_2 .

XNOR Gate

x_1	x_2	y	a_1	a_2
0	0	1	0.002	0.881
0	1	0	0.119	0.119
1	0	0		
1	1	1		

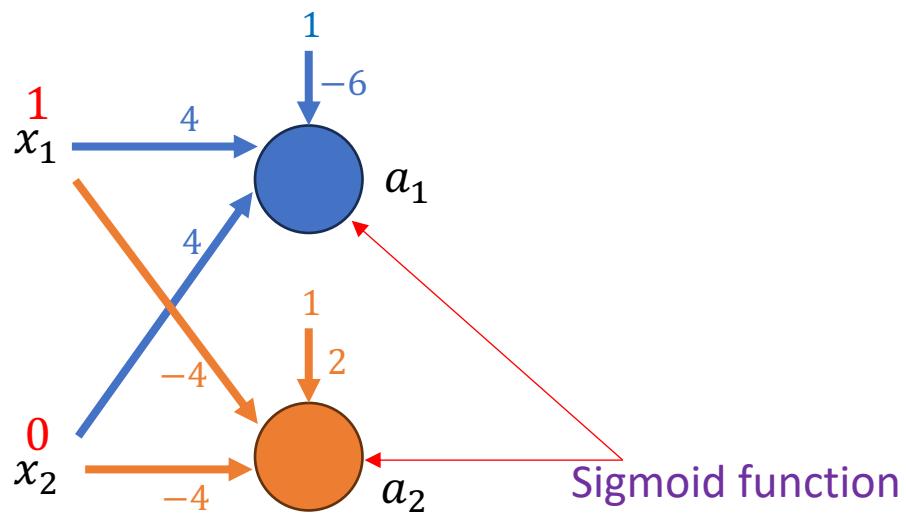
$$\begin{aligned}a_1 &= \sigma(-6 + 4x_1 + 4x_2) \\&= \sigma(-6 + 4 \times 0 + 4 \times 1) = 0.119\end{aligned}$$

$$\begin{aligned}a_2 &= \sigma(2 + (-4)x_1 + (-4)x_2) \\&= \sigma(2 + (-4) \times 0 + (-4) \times 1) = 0.119\end{aligned}$$

XNOR Gate Modelling

With More Neurons

Neurons can work as learnable feature transformation function.



\$a_1\$ and \$a_2\$ are new features created based on existing features \$x_1\$ and \$x_2\$.

XNOR Gate

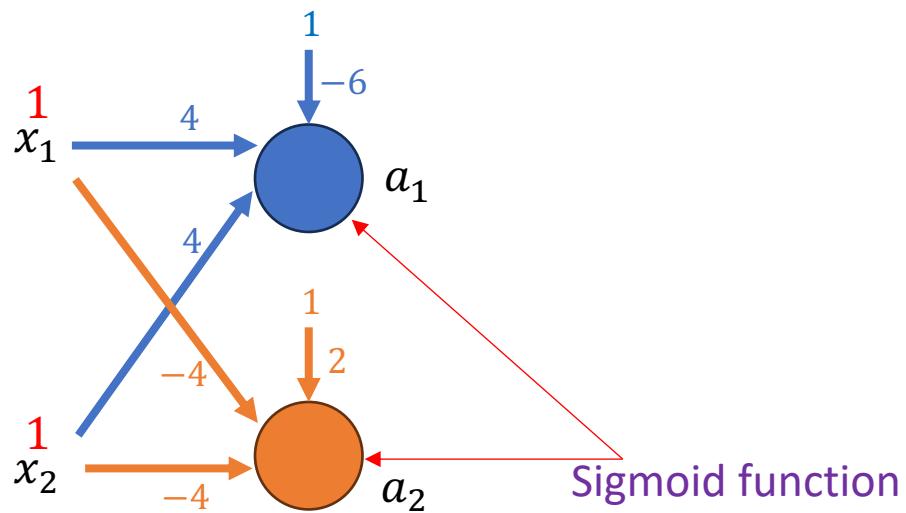
\$x_1\$	\$x_2\$	\$y\$	\$a_1\$	\$a_2\$
0	0	1	0.002	0.881
0	1	0	0.119	0.119
1	0	0	0.119	0.119
1	1	1		

$$\begin{aligned}a_1 &= \sigma(-6 + 4x_1 + 4x_2) \\&= \sigma(-6 + 4 \times 1 + 4 \times 0) = 0.119 \\ \\a_2 &= \sigma(2 + (-4)x_1 + (-4)x_2) \\&= \sigma(2 + (-4) \times 1 + (-4) \times 0) = 0.119\end{aligned}$$

XNOR Gate Modelling

With More Neurons

Neurons can work as learnable feature transformation function.



\$a_1\$ and \$a_2\$ are new features created based on existing features \$x_1\$ and \$x_2\$.

XNOR Gate

\$x_1\$	\$x_2\$	\$y\$	\$a_1\$	\$a_2\$
0	0	1	0.002	0.881
0	1	0	0.119	0.119
1	0	0	0.119	0.119
1	1	1	0.881	0.002

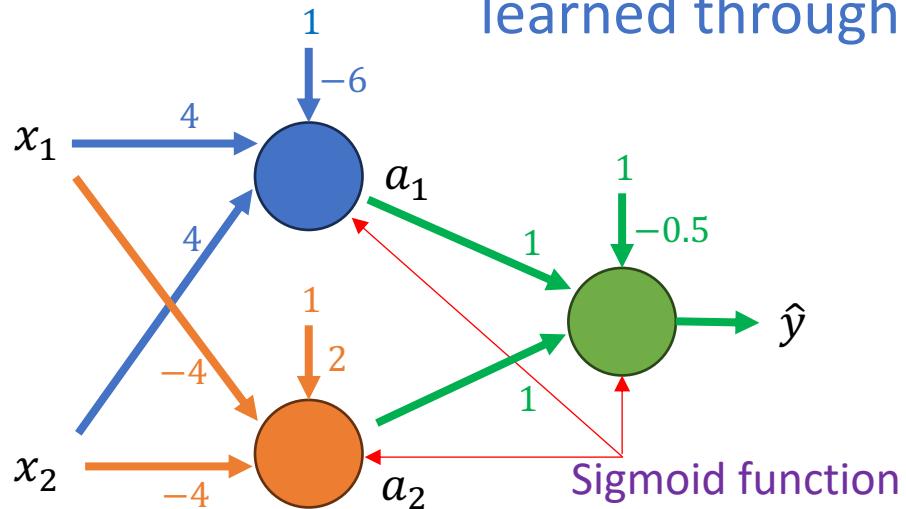
$$\begin{aligned}a_1 &= \sigma(-6 + 4x_1 + 4x_2) \\&= \sigma(-6 + 4 \times 1 + 4 \times 1) = 0.881 \\a_2 &= \sigma(2 + (-4)x_1 + (-4)x_2) \\&= \sigma(2 + (-4) \times 1 + (-4) \times 1) = 0.002\end{aligned}$$

XNOR Gate Modelling

With More Neurons

Neurons can work as learnable feature transformation function. **Why learnable?**

Weights in neurons are learned through training.

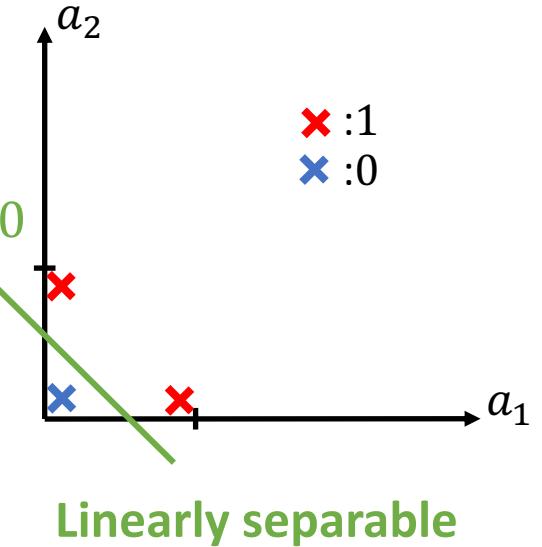


a_1 and a_2 are new features created based on existing features x_1 and x_2 .

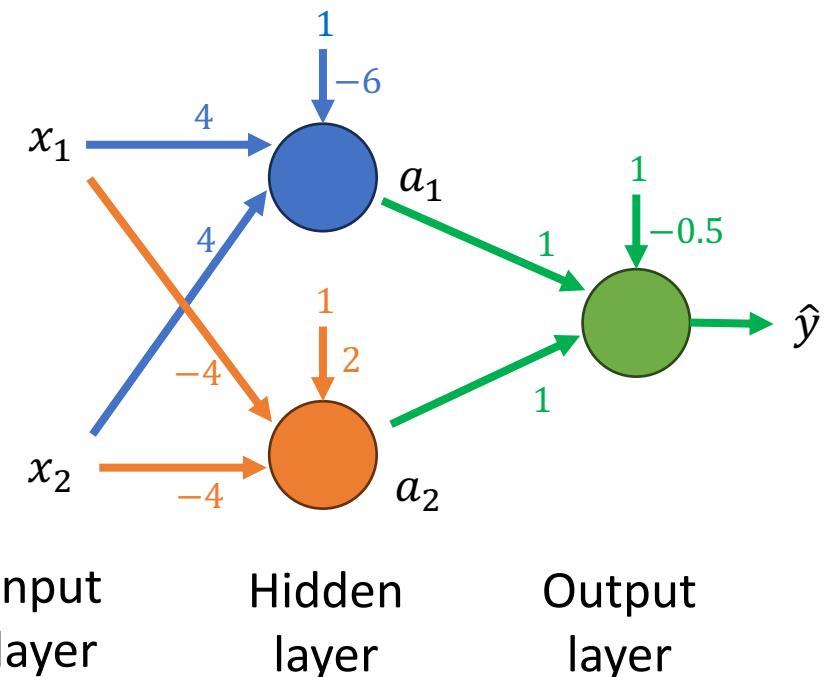
$$z = -0.5 + x_1 + x_2 = 0$$

XNOR Gate		
x_1	x_2	y
0	0	1
0	1	0
1	0	0
1	1	1

a_1	a_2
0.002	0.881
0.119	0.119
0.119	0.119
0.881	0.002



Multi-layer Neural Network

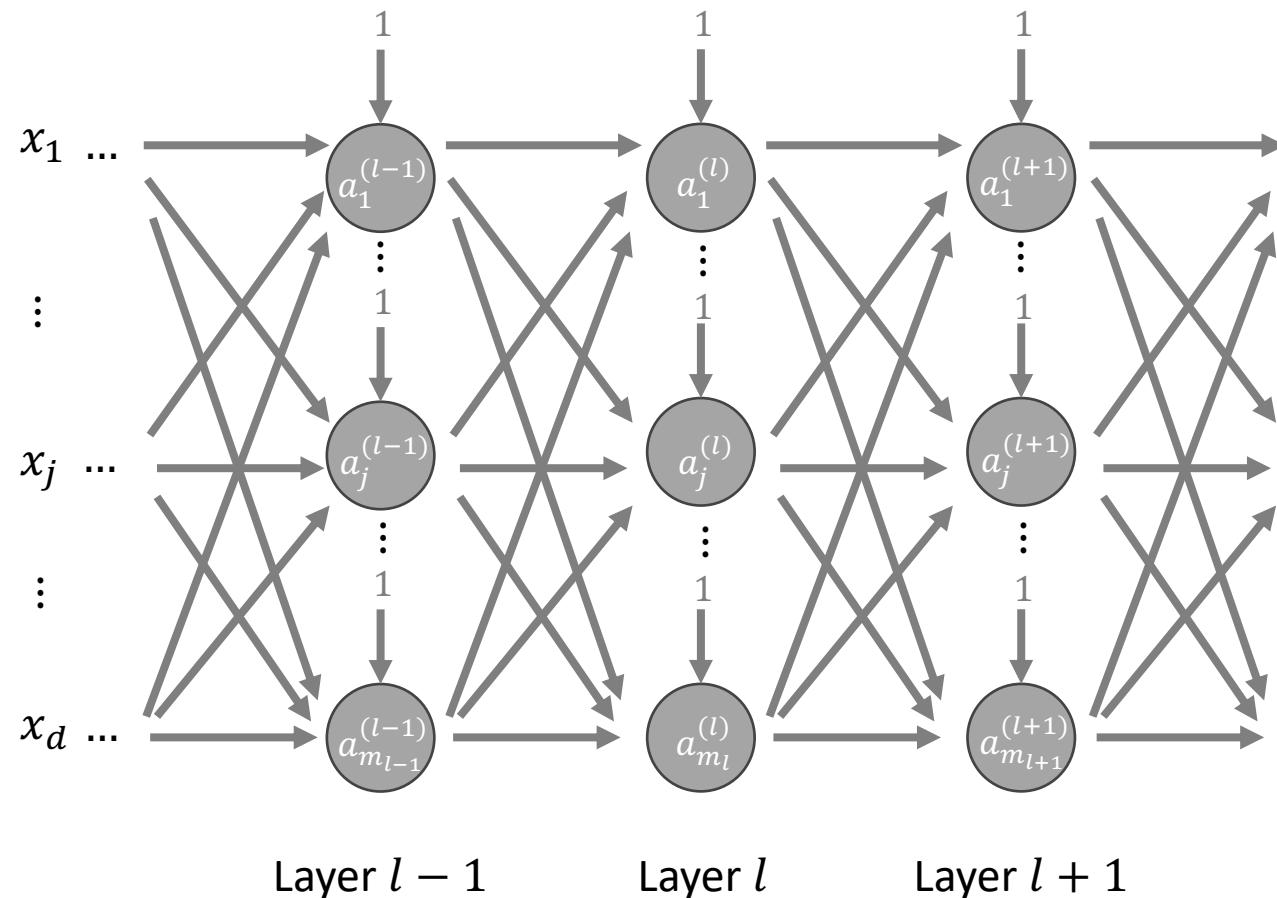


A **hidden layer** is any layer between the input and output layers.

Both the hidden layer and output layer are **fully-connected layers**. Fully-connected layer is a type of layer in a neural network where every neuron is connected to every feature/neuron in the previous layer.

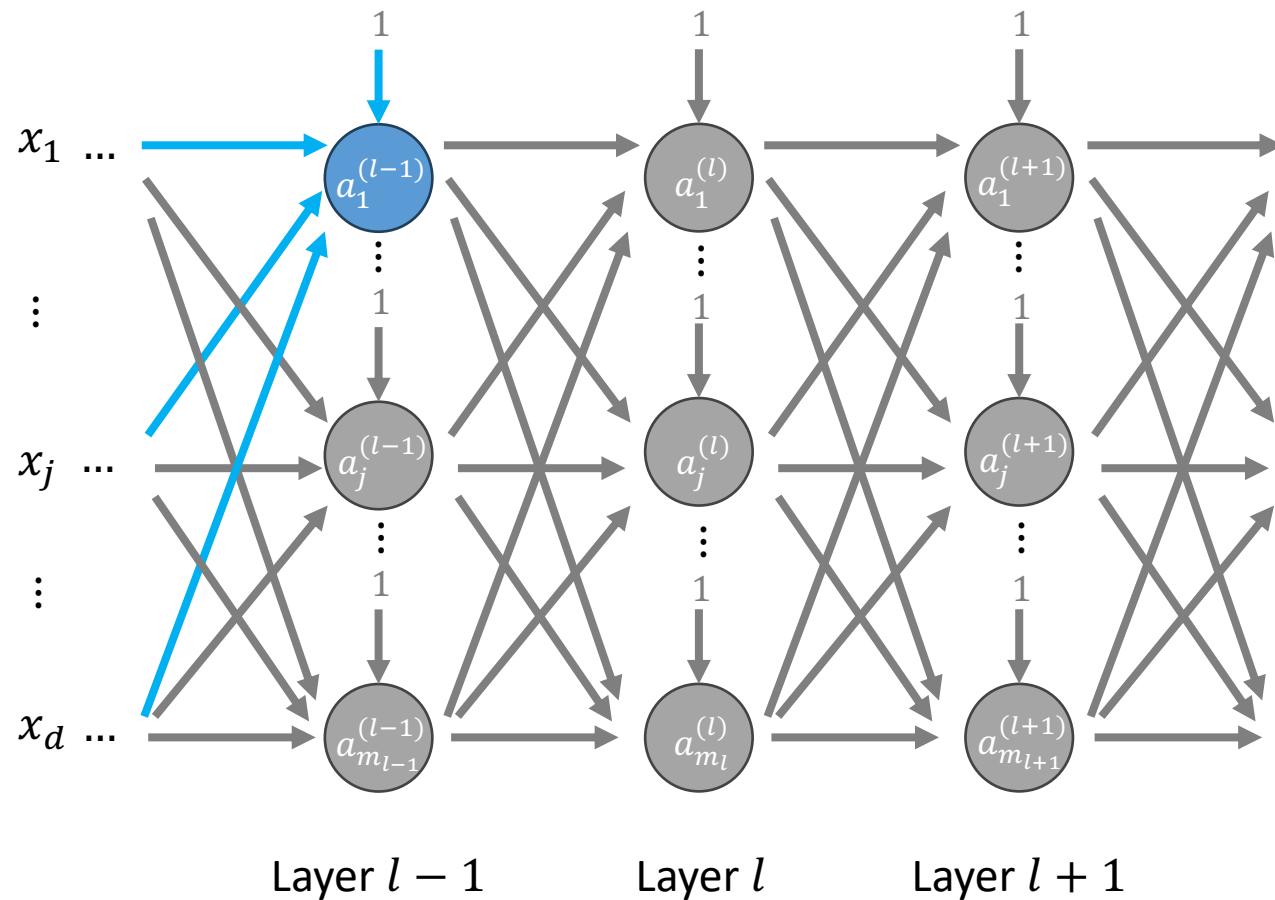
Multi-layer Neural Network with only fully-connected layers can also be called **Multi-layer Perceptron**.

Multi-layer Neural Network



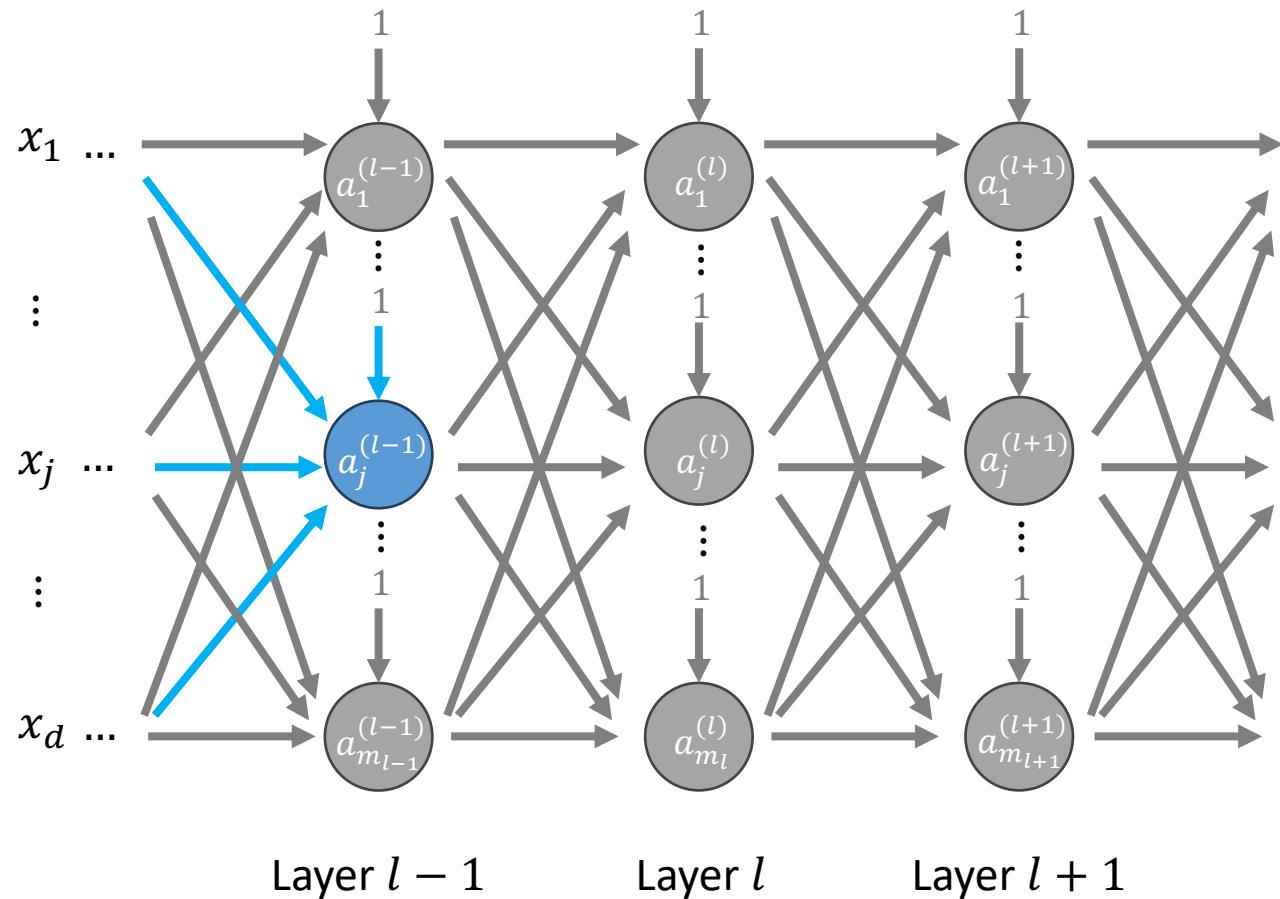
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



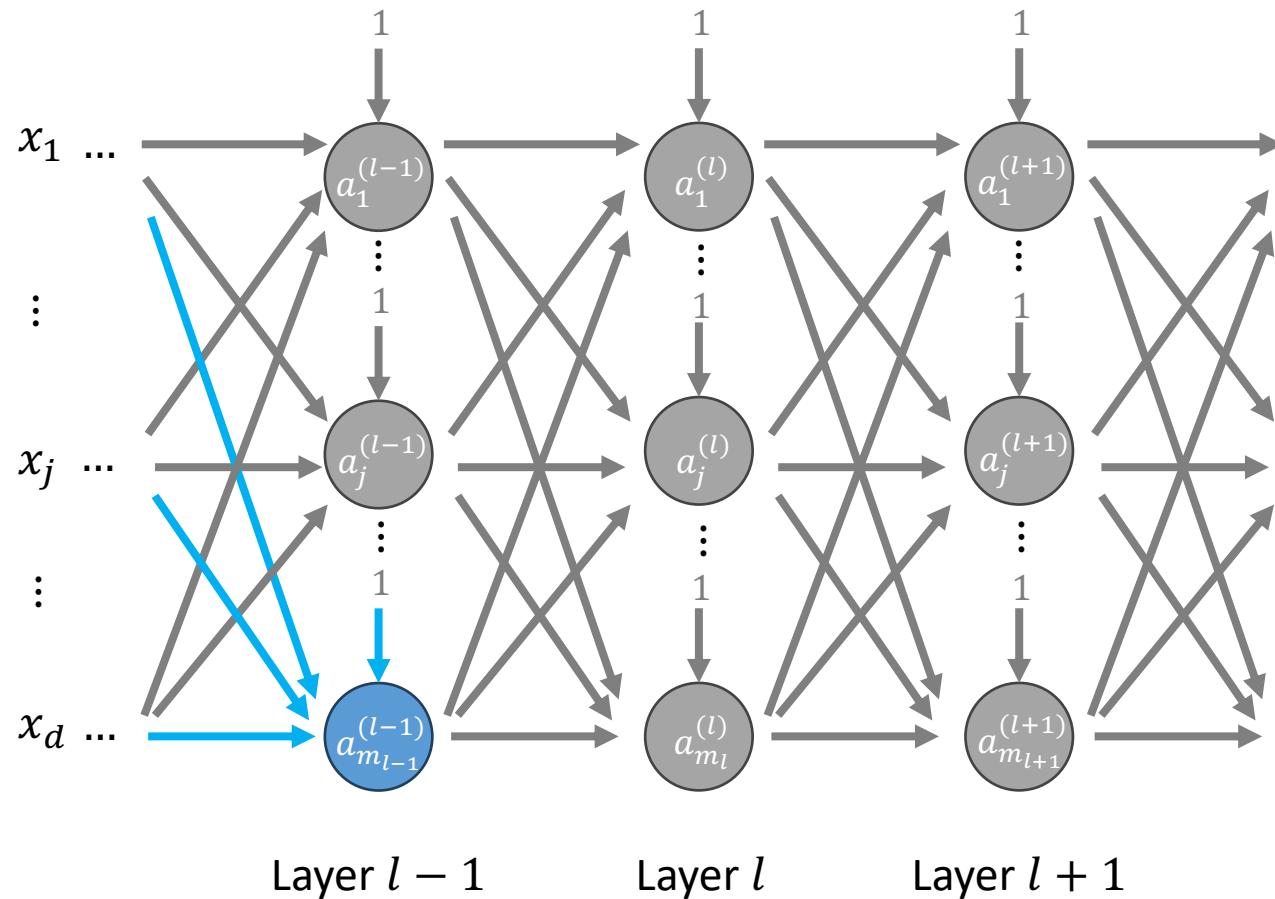
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



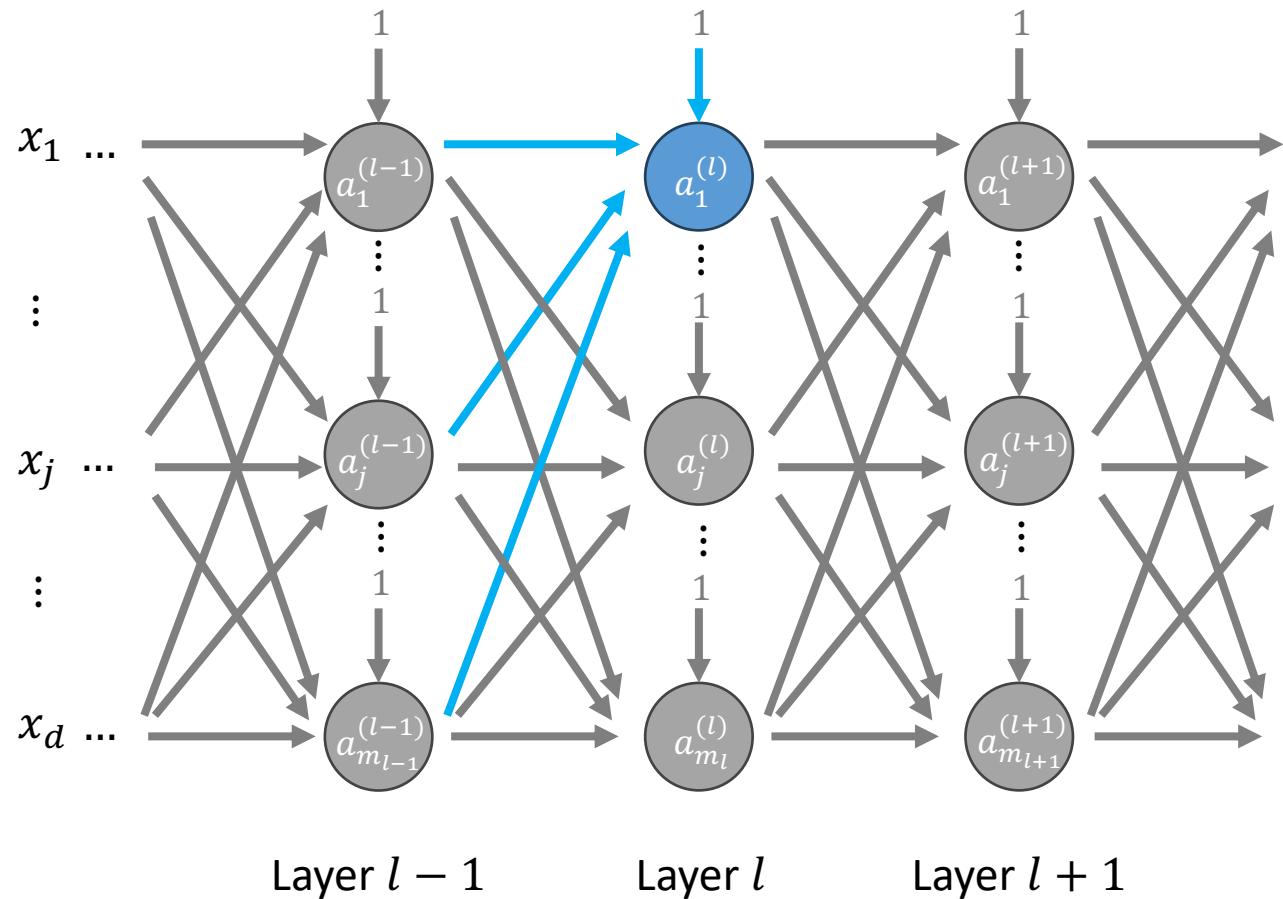
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



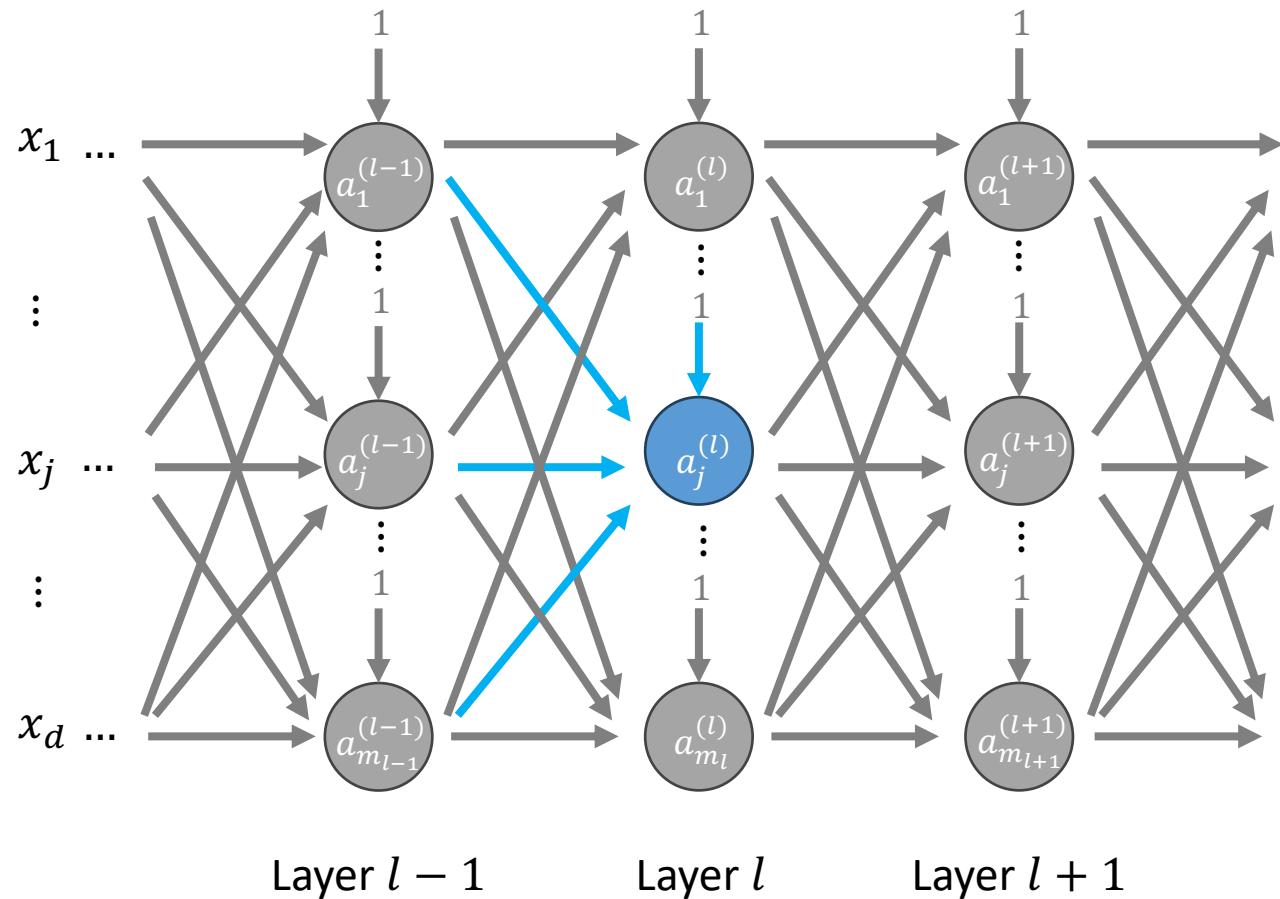
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



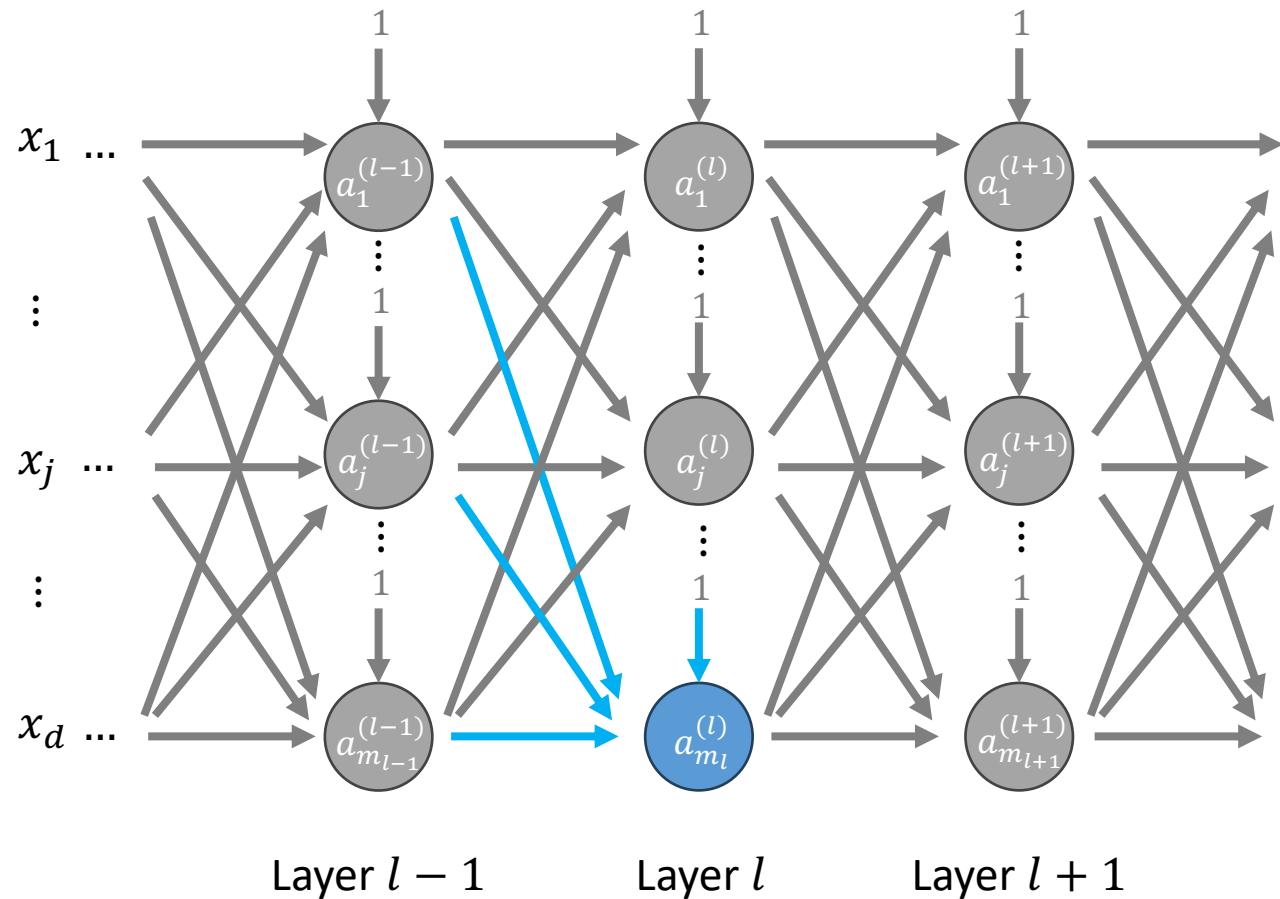
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



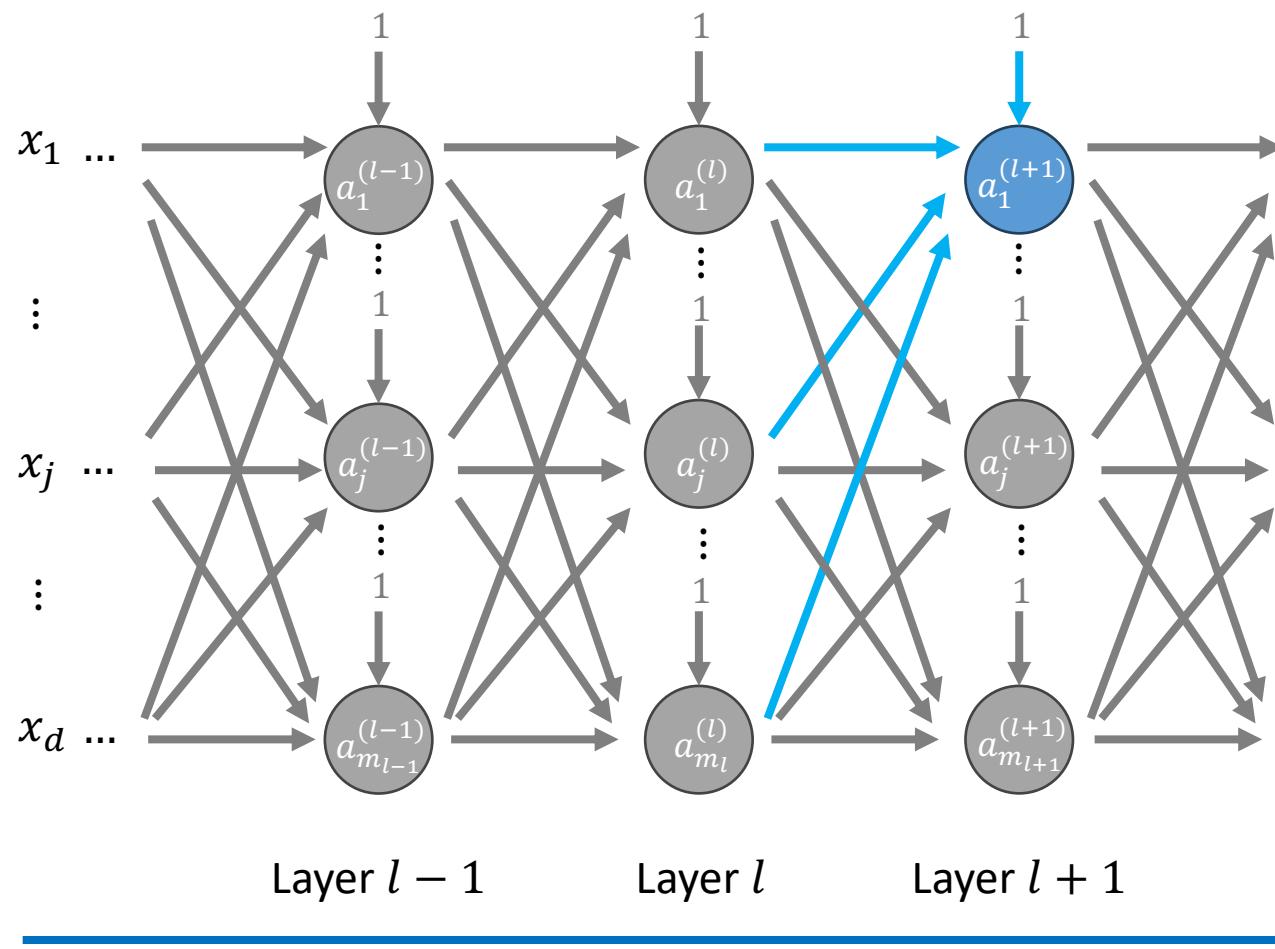
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Multi-layer Neural Network



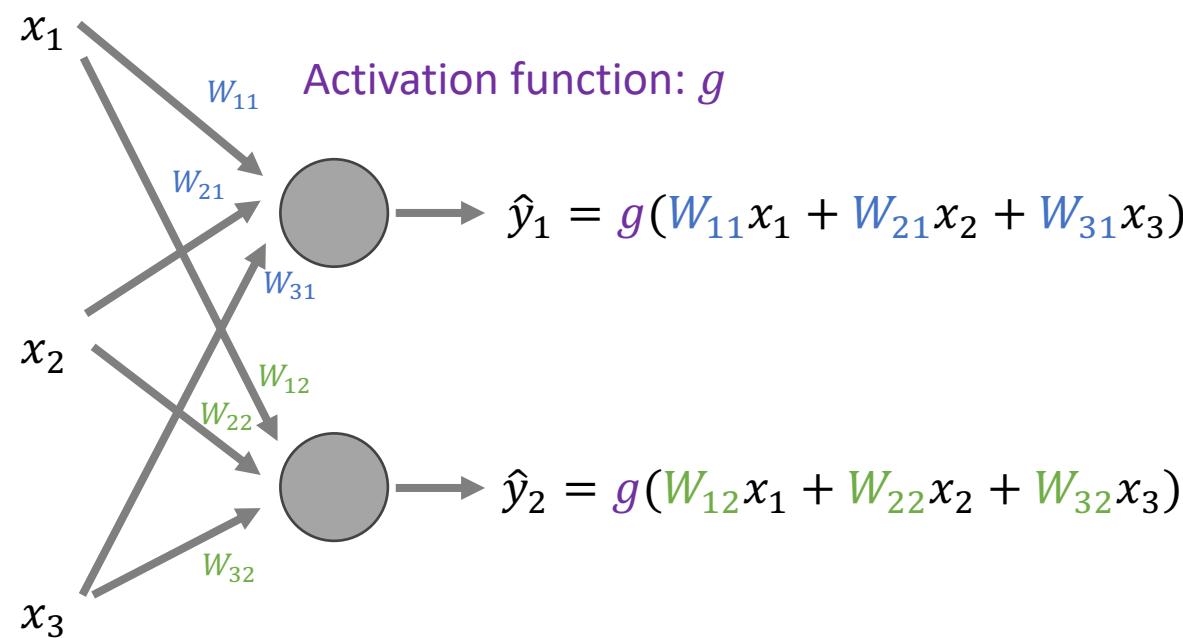
Layer $l - 1$, Layer l and Layer $l + 1$
are **fully-connected layers**

Forward Propagation

The process by which input data passes through a neural network to generate the output.

Neural Network and Matrix Multiplication

Single-layer (No bias)



Input (number of weights per neuron / input variables)

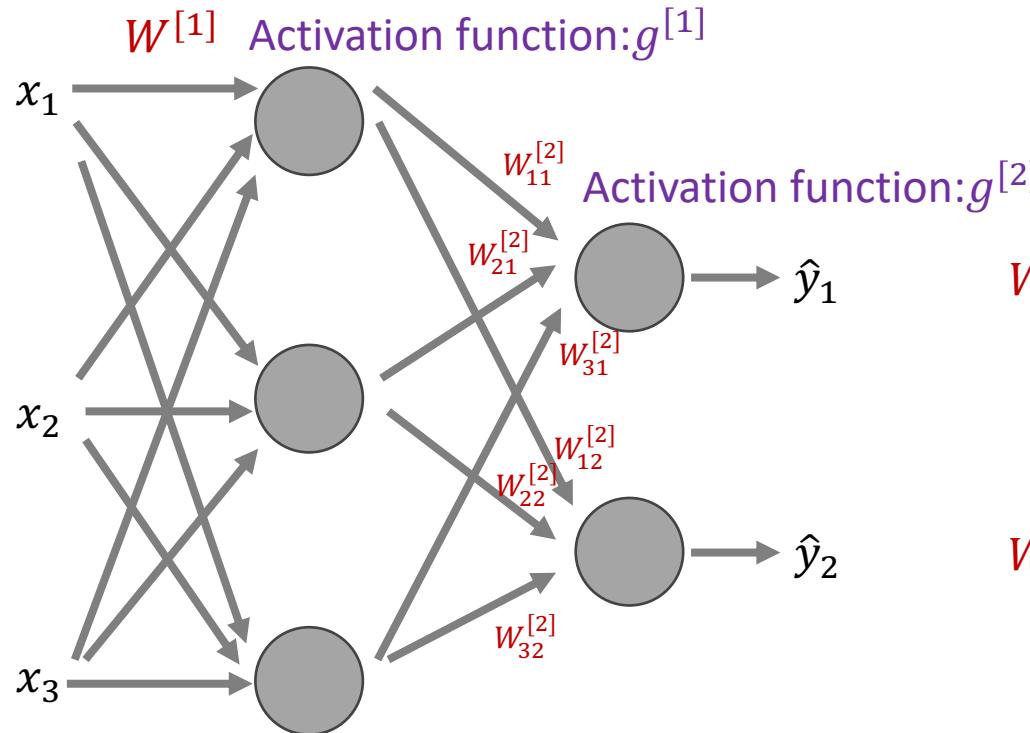
$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} W_{11} & W_{12} \\ W_{21} & W_{22} \\ W_{31} & W_{32} \end{bmatrix}$$
$$\hat{\mathbf{y}} = g(\mathbf{W}^T \mathbf{x}) = g\left(\begin{bmatrix} W_{11} & W_{12} \\ W_{21} & W_{22} \\ W_{31} & W_{32} \end{bmatrix}^T \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = g\left(\begin{bmatrix} W_{11} & W_{21} & W_{31} \\ W_{12} & W_{22} & W_{32} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \end{bmatrix}$$

Output (number of layer's neurons / output variables)

Neural Network and Matrix Multiplication

Multi-layer (No bias)

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$



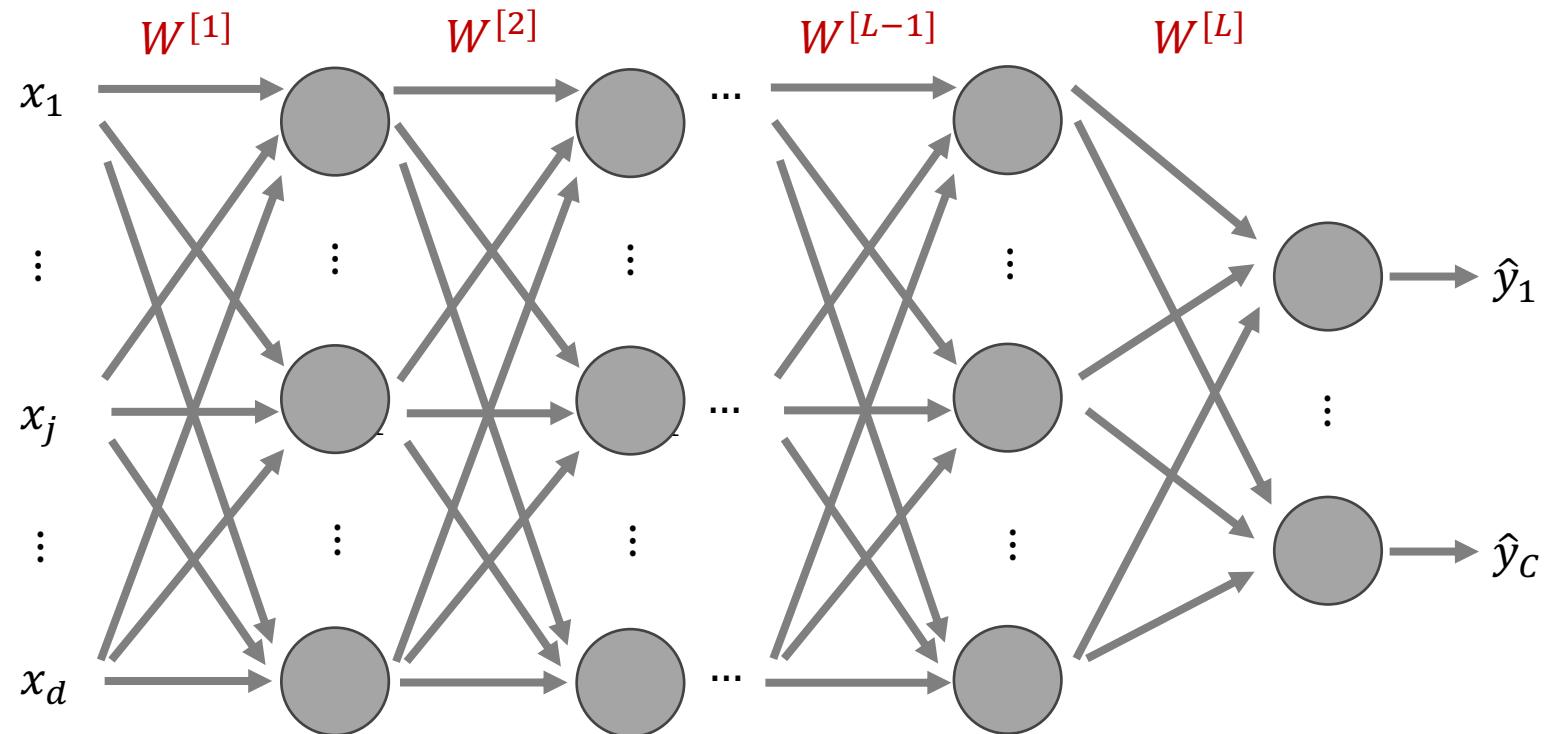
$$W^{[1]} = \begin{bmatrix} W_{11}^{[1]} & W_{12}^{[1]} & W_{13}^{[1]} \\ W_{21}^{[1]} & W_{22}^{[1]} & W_{23}^{[1]} \\ W_{31}^{[1]} & W_{32}^{[1]} & W_{33}^{[1]} \end{bmatrix}$$

$$W^{[2]} = \begin{bmatrix} W_{11}^{[2]} & W_{12}^{[2]} \\ W_{21}^{[2]} & W_{22}^{[2]} \\ W_{31}^{[2]} & W_{32}^{[2]} \end{bmatrix}$$

$$\hat{\mathbf{y}} = g^{[2]} \left(W^{[2]^T} g^{[1]} \left(W^{[1]^T} \mathbf{x} \right) \right) = g^{[2]} \left(\begin{bmatrix} W_{11}^{[2]} & W_{12}^{[2]} \\ W_{21}^{[2]} & W_{22}^{[2]} \\ W_{31}^{[2]} & W_{32}^{[2]} \end{bmatrix}^T g^{[1]} \left(\begin{bmatrix} W_{11}^{[1]} & W_{12}^{[1]} & W_{13}^{[1]} \\ W_{21}^{[1]} & W_{22}^{[1]} & W_{23}^{[1]} \\ W_{31}^{[1]} & W_{32}^{[1]} & W_{33}^{[1]} \end{bmatrix}^T \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \right) \right) = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \end{bmatrix}$$

Neural Network and Matrix Multiplication

Multi-layer (No bias)



$$\hat{\mathbf{y}} = g^{[L]} \left(W^{[L]^T} \dots g^{[L-1]} \left(W^{[L-1]^T} \dots g^{[l]} \left(W^{[l]^T} \dots g^{[2]} \left(W^{[2]^T} g^{[1]} \left(W^{[1]^T} \mathbf{x} \right) \right) \right) \right) \right) = \begin{bmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_c \end{bmatrix}$$

Outline

- Neural Network
 - Neuron
 - Activation Functions
 - Multi-layer Neural Network
- **Tasks for Neural Network**
 - Binary Classification
 - Multi-class Classification
 - Single-output Regression
 - Multi-output Regression
- Backpropagation

NN for Binary Classification - Data

Suppose:

- We are given N data points.
- Each data point consists of **features** and a **target** variable.
- The features are described by a vector of **real numbers** in dimension d .
- The target is $\{0,1\}$, where 0 is “negative” class and 1 is “positive” class.

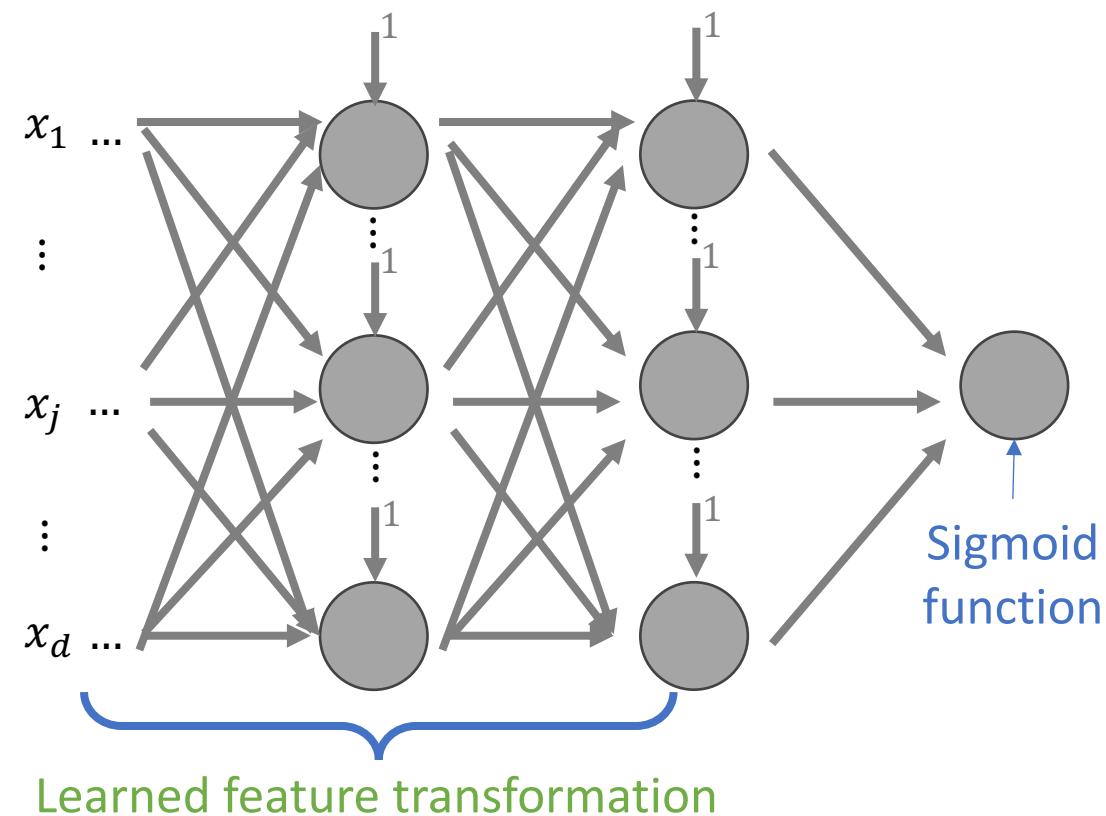
NN for Binary Classification - Model

Hidden layers

- Learn feature transformation

Output layer:

- Number of neurons: 1
- Activation function: Sigmoid



NN for Multi-Class Classification - Data

Suppose:

- We are given N data points.
- Each data point consists of **features** and a **target label**.
- The features are described by a vector of **real numbers** in dimension d .
- The target is the class label $\in \{\text{class 1, ..., class } C\}$. C is the total number of classes

One-hot Encoding

One-hot encoding is a way to represent categorical labels as vectors of numbers that a neural network can understand.

Each class is represented by a vector that has:

- 1 in the position corresponding to that class, and
- 0s everywhere else
- **Length of vector == The number of classes**

Target Label	Target Vector
Cat	[1, 0, 0]
Dog	[0, 1, 0]
Bird	[0, 0, 1]

NN for Multi-Class Classification - Model

Hidden layers

- Learn feature transformation

Output layer

- Number of neurons: C

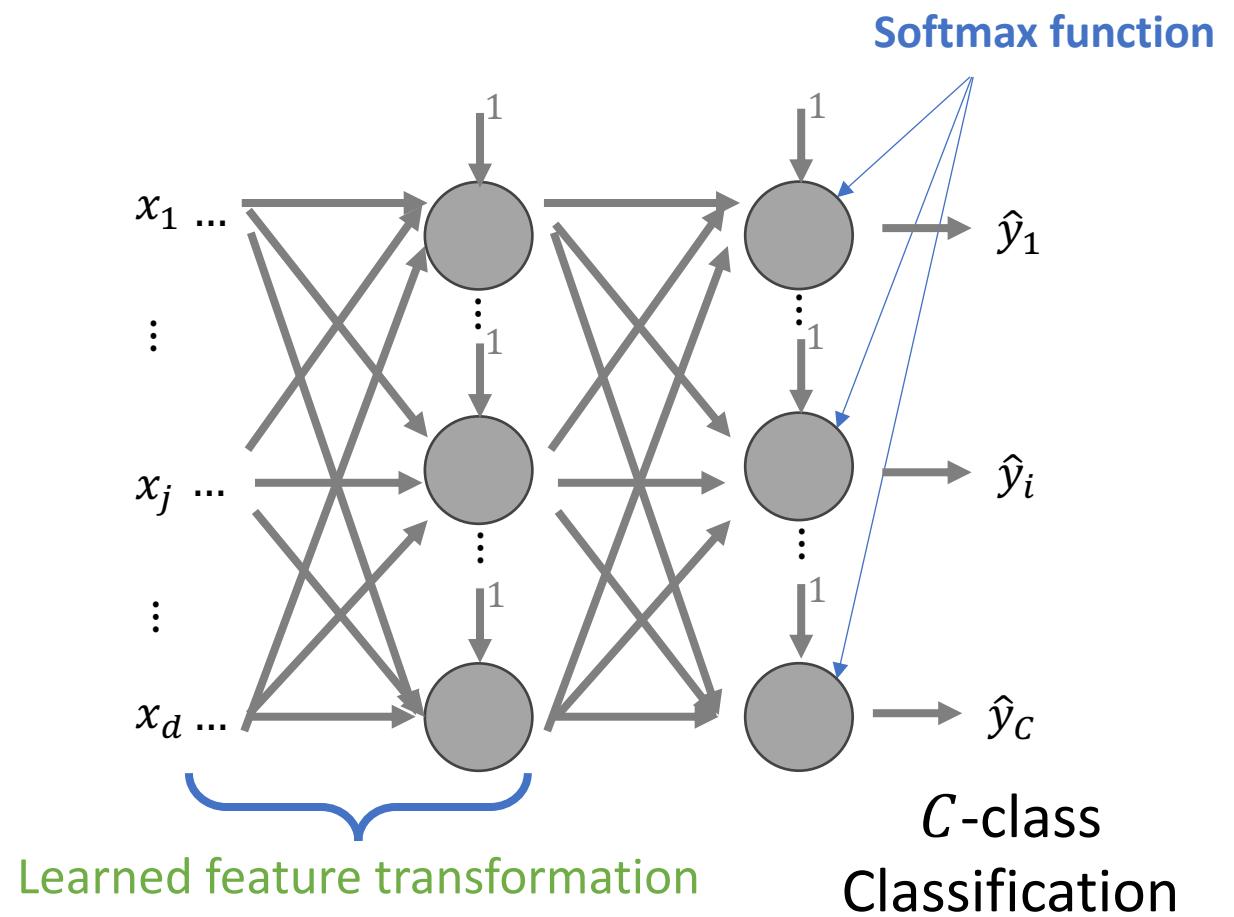
One data point $x = [x_1, x_2, \dots, x_d]^T$

Target vector $y = [y_1, y_2, \dots, y_C]^T$

Predict output $\hat{y} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_C]^T$

- Pick the class with the **highest** \hat{y}_i

- Activation function: Softmax

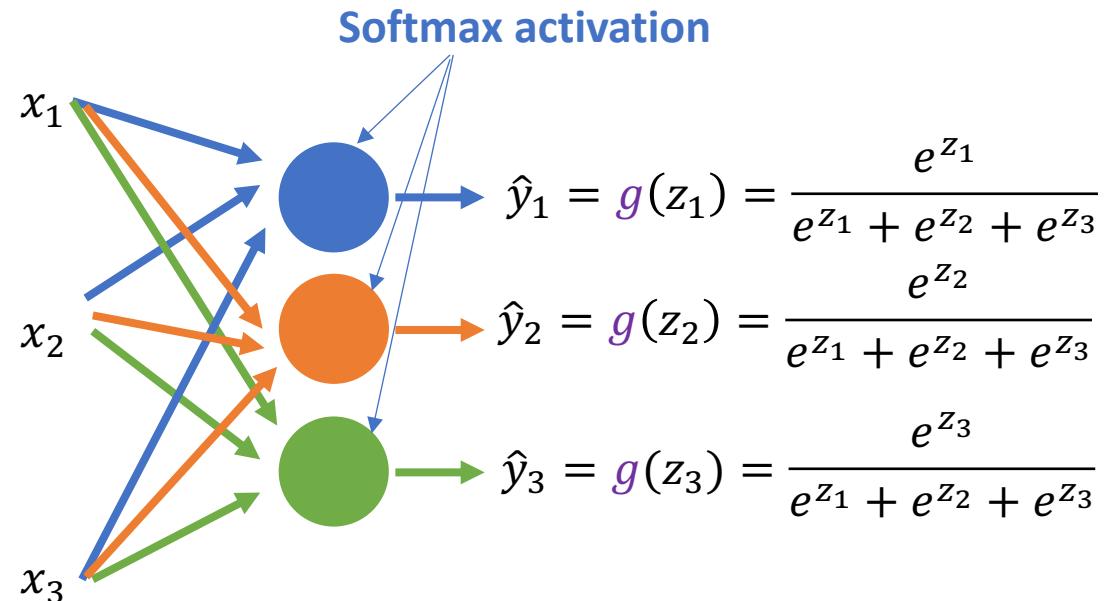


Softmax Function

Softmax function computes the probability of belonging to class i as:

$$g(z_i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}} \in [0,1]$$

Here, z_i is the weighted sum value computed in neuron i .



NN for Single-Output Regression - Data

Suppose:

- We are given N data points.
- Each data point consists of **features** and a **target** variable.
- Each data point has d features and they are described by **real numbers**.
- The target is also a **real number**.

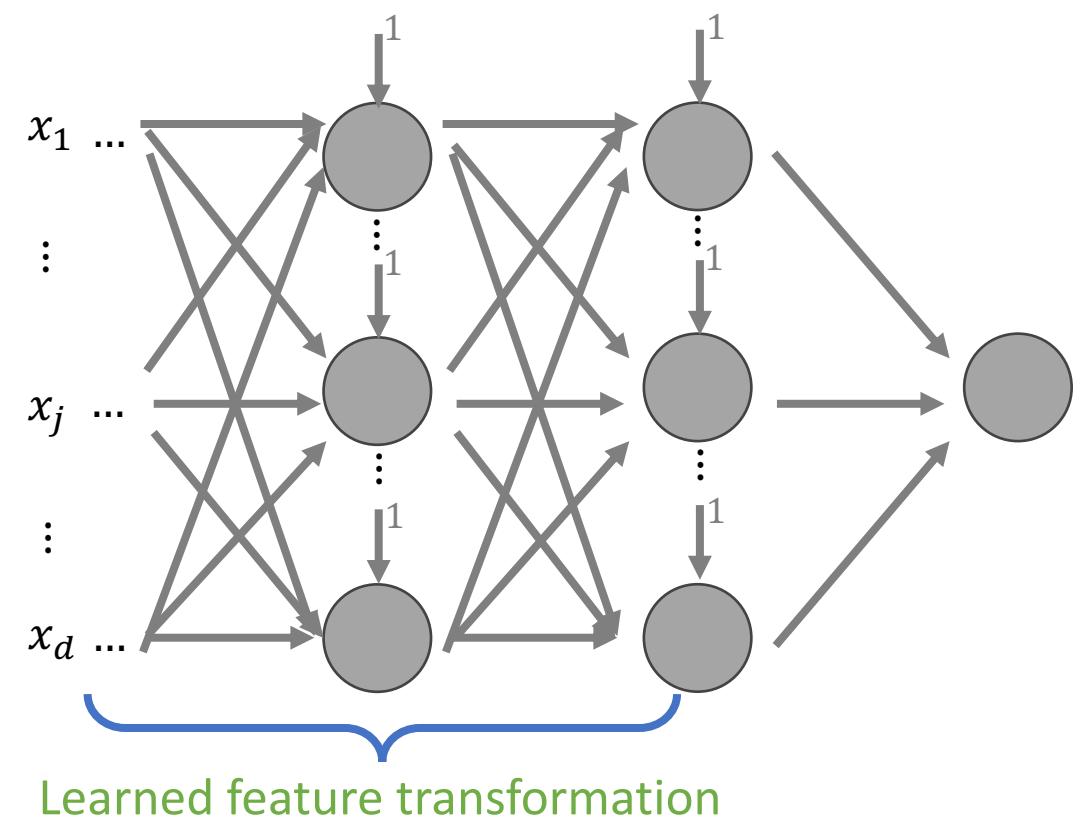
NN for Single-Output Regression - Model

Hidden layers

- Learn feature transformation.

Output layer

- Number of neurons: 1
- Activation function:
 - Identity/No activation function
 - Sigmoid
 - Tanh
 - ...



NN for Multi-Output Regression - Data

Suppose:

- We are given N data points.
- Each data point consists of **features** and **target vectors**.
- Each data point has d features and they are described by **real numbers**.
- The target for each data point is a vector containing K real numbers.

NN for Multi-Output Regression - Model

Hidden layers

- Learn feature transformation.

Output layer

- Number of neurons: K

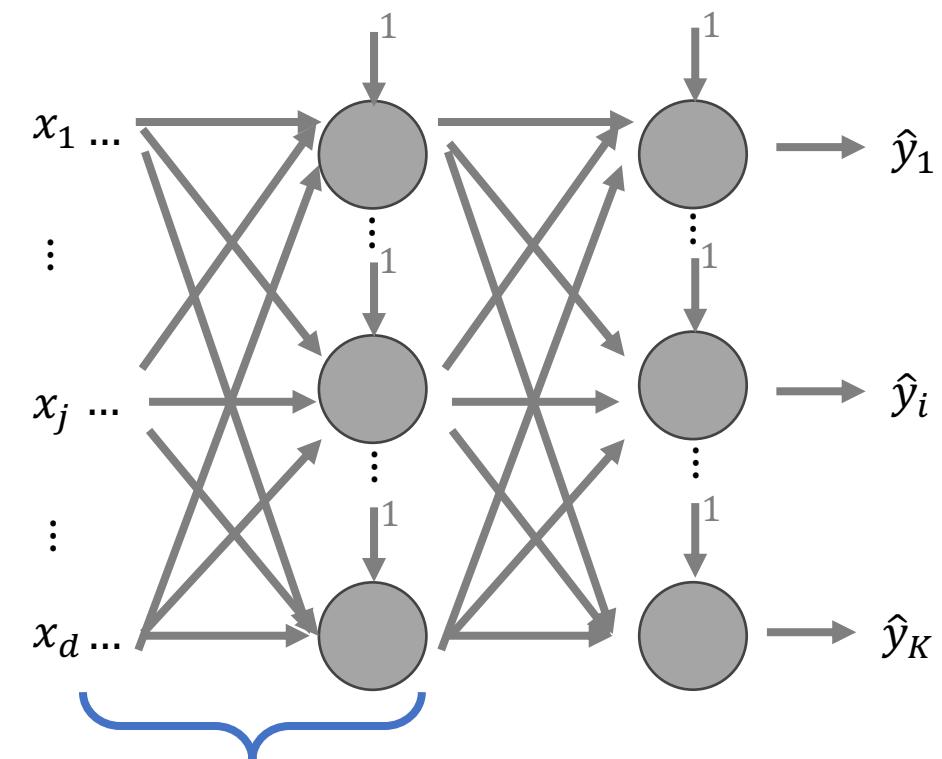
One data point $\mathbf{x} = [x_1, x_2, \dots, x_d]^T$

Target vector $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$

Predict output $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$

- Activation function:

- Identity/No activation function
- Sigmoid
- Tanh
- ...



Learned feature transformation

Outline

- Neural Network
 - Neuron
 - Activation Functions
 - Multi-layer Neural Network
- Tasks for Neural Network
 - Binary Classification
 - Multi-class Classification
 - Single-output Regression
 - Multi-output Regression
- **Backpropagation**

Learning via Gradient Descent

Model Weights

- Step1: Start at some \mathbf{w} (e.g., randomly initialized).
- Step2: Update \mathbf{w} a step to the opposite direction of the gradient (i.e., towards lower loss)

$$\mathbf{w}_j \leftarrow \mathbf{w}_j - \gamma \frac{\partial J(\mathbf{w}_0, \mathbf{w}_1, \dots)}{\partial \mathbf{w}_j}.$$

Loss function
Learning Rate

- Repeat Step 2 until termination criterion is satisfied.
 - E.g., change between steps is small, maximum number of steps is reached, etc

The gradient descent can also be used to update the weights in neural network.

Background: Chain Rule (1)

- The chain rule is a formula in calculus used to compute the derivative of a composition of functions, e.g.,:

$$l = h(g(f(x)))$$

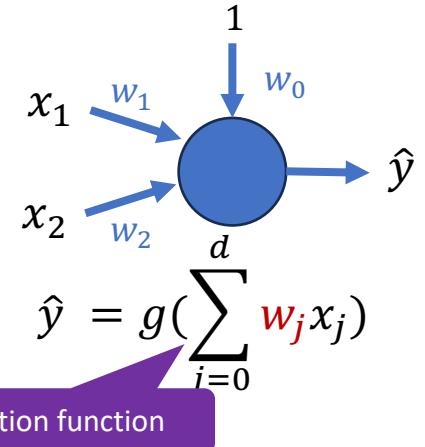
- By introducing the intermediate variables, we can rewrite the composition as:
Let $z = f(x)$ and $y = g(z)$, then $l = h(y)$
- The chain rule states that the derivative of l with respect to x is given by

$$\Delta x \rightarrow \Delta z \rightarrow \Delta y \rightarrow \Delta l \quad \frac{dl}{dx} = \frac{dl}{dy} \frac{dy}{dz} \frac{dz}{dx}$$

Gradient Computation

1. Generate the predicted value \hat{y} for a given data point (x, y) :

$$\hat{y} = g\left(\sum_{j=0}^d w_j x_j\right)$$



Activation function

2. Compute the loss, e.g., MSE:

$$L = \frac{1}{2}(\hat{y} - y)^2$$

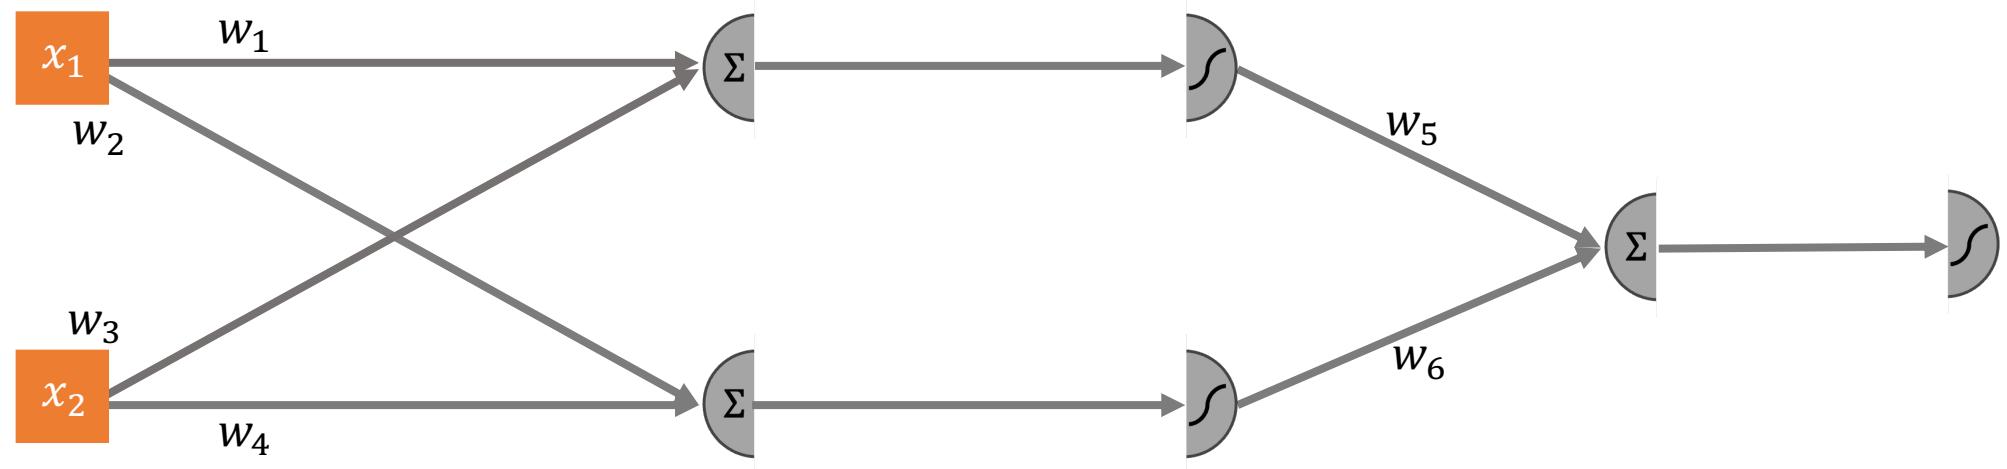
3. Let $z = \sum_{j=0}^d w_j x_j$ and $\hat{y} = g(z)$, the gradient of loss function with respect to w_j is:

$$\frac{\partial L}{\partial w_j} = \frac{dL}{d\hat{y}} \frac{d\hat{y}}{dz} \frac{\partial z}{\partial w_j}$$

$$\frac{dL}{d\hat{y}} = \frac{d\left(\frac{1}{2}(\hat{y} - y)^2\right)}{d\hat{y}} = (\hat{y} - y) \quad \frac{d\hat{y}}{dz} = \frac{d(g(z))}{dz} = g'(z) \quad \frac{\partial z}{\partial w_j} = \frac{\partial(\sum_{j=0}^d w_j x_j)}{\partial w_j} = \frac{\partial(w_j x_j + \sum_{k \neq j} w_k x_k)}{\partial w_j} = x_j$$

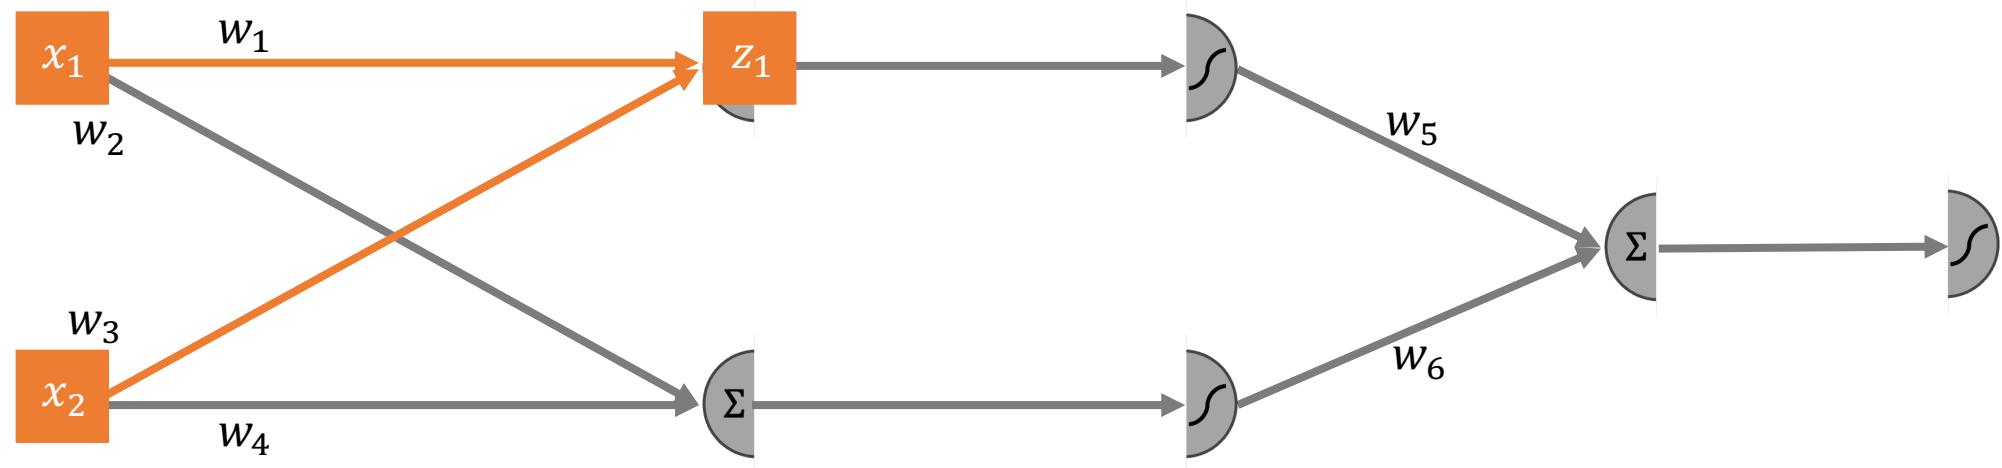
Gradient Computation for Multi-layer NN

1. Generate the predicted value \hat{y} for a given data point (x, y) :



Gradient Computation for Multi-layer NN

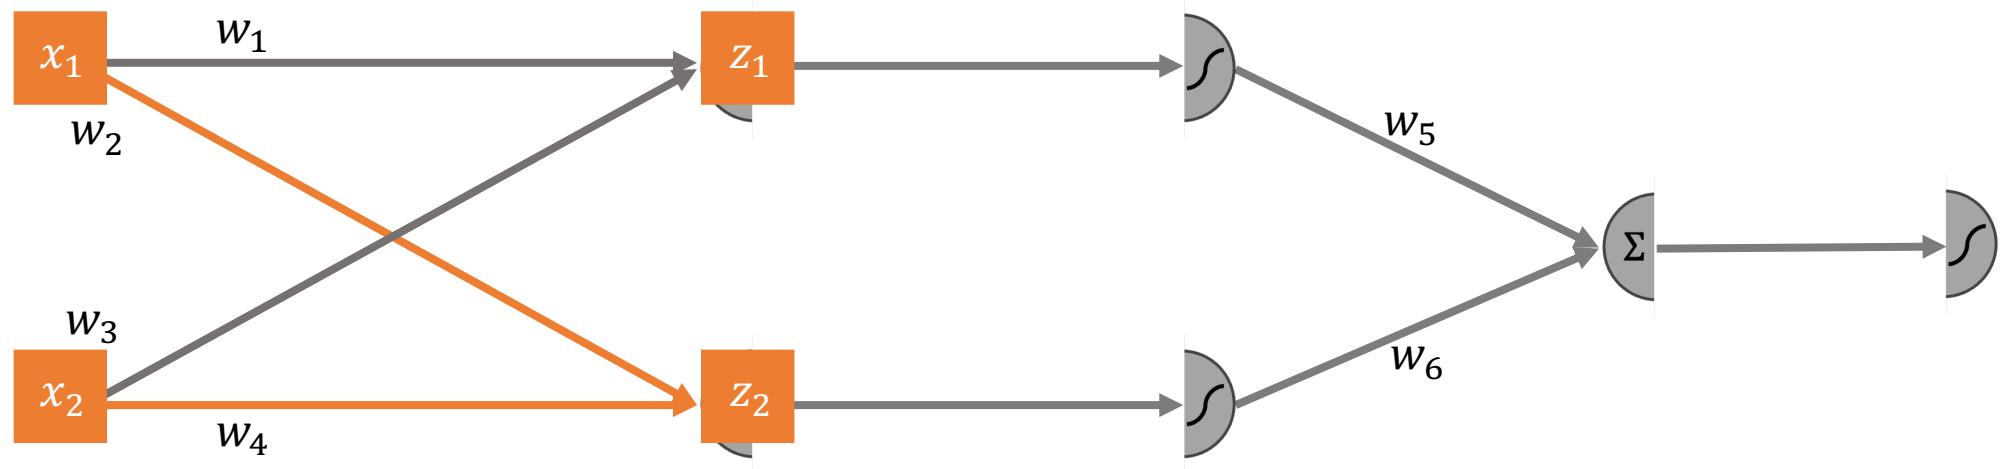
1. Generate the predicted value \hat{y} for a given data point (x, y) :



$$z_1 = w_1 x_1 + w_3 x_2$$

Gradient Computation for Multi-layer NN

1. Generate the predicted value \hat{y} for a given data point (x, y) :



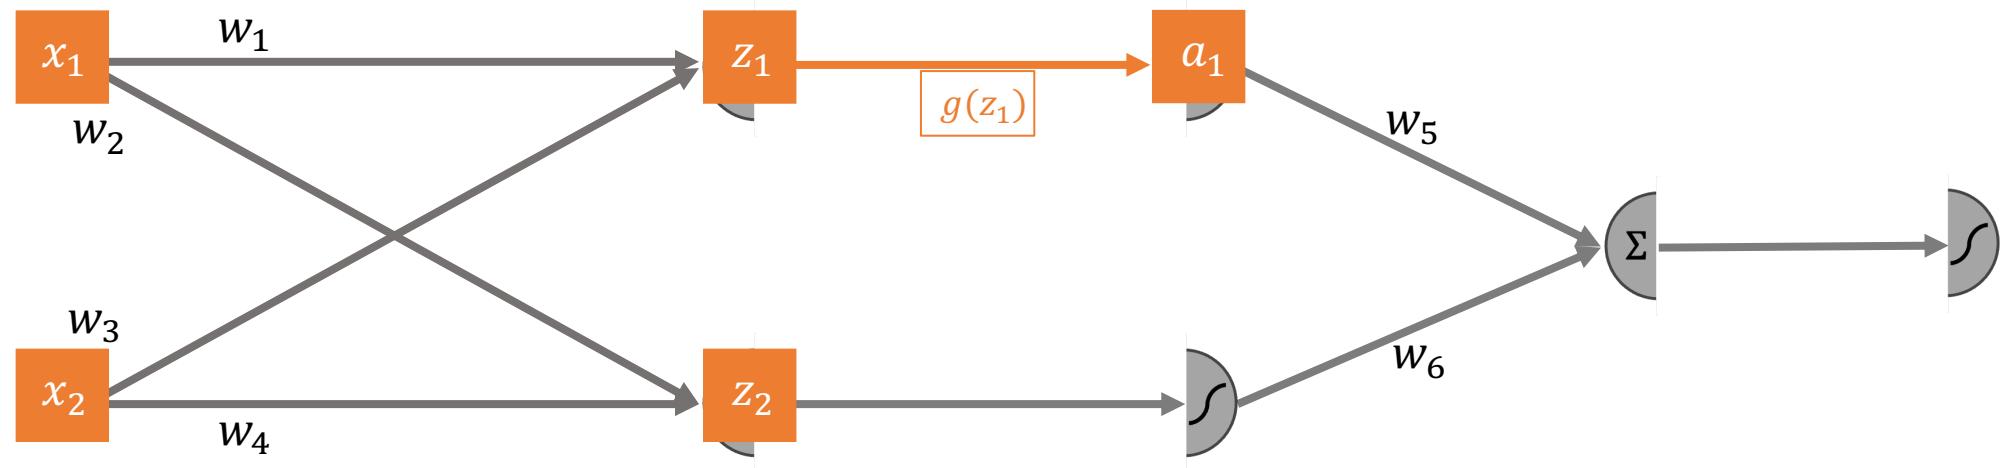
$$z_1 = w_1 x_1 + w_3 x_2$$

$$z_2 = w_2 x_1 + w_4 x_2$$

Gradient Computation for Multi-layer NN

Activation function: $g(z)$

1. Generate the predicted value \hat{y} for a given data point (x, y) :



$$z_1 = w_1 x_1 + w_3 x_2$$

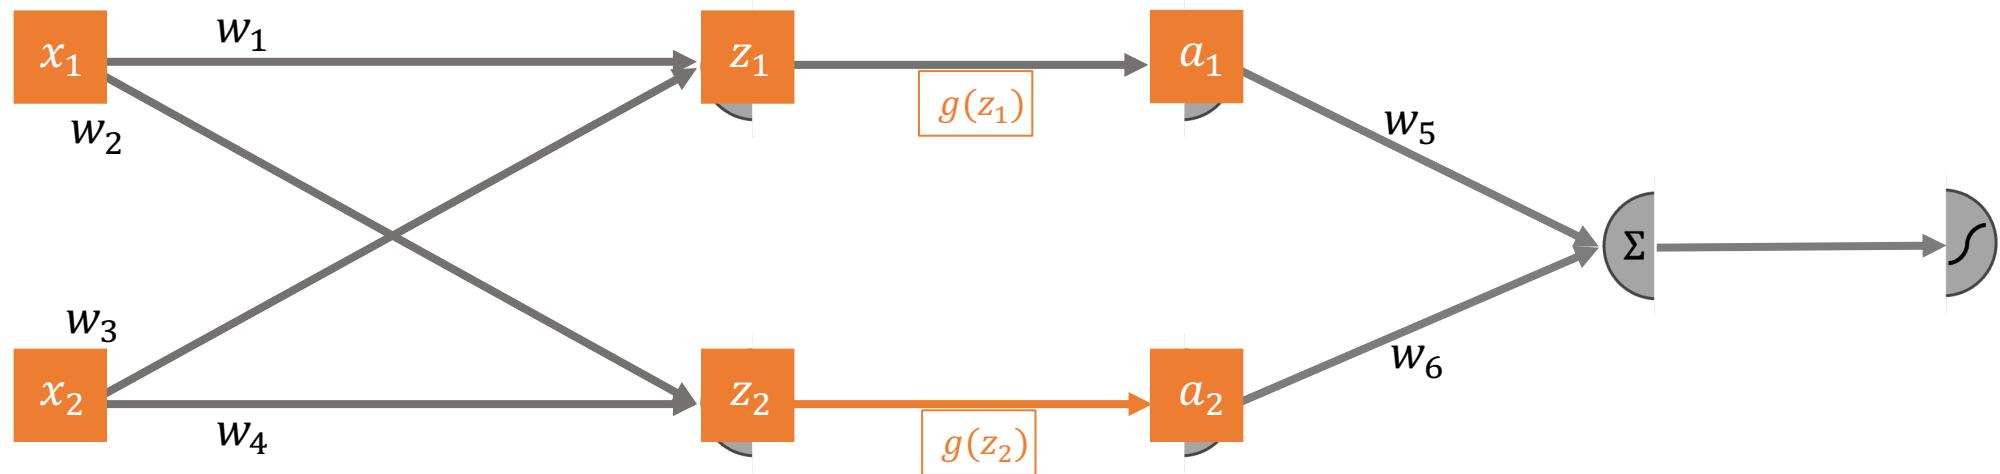
$$z_2 = w_2 x_1 + w_4 x_2$$

$$a_1 = g(z_1)$$

Gradient Computation for Multi-layer NN

Activation function: $g(z)$

1. Generate the predicted value \hat{y} for a given data point (x, y) :



$$z_1 = w_1 x_1 + w_3 x_2$$

$$z_2 = w_2 x_1 + w_4 x_2$$

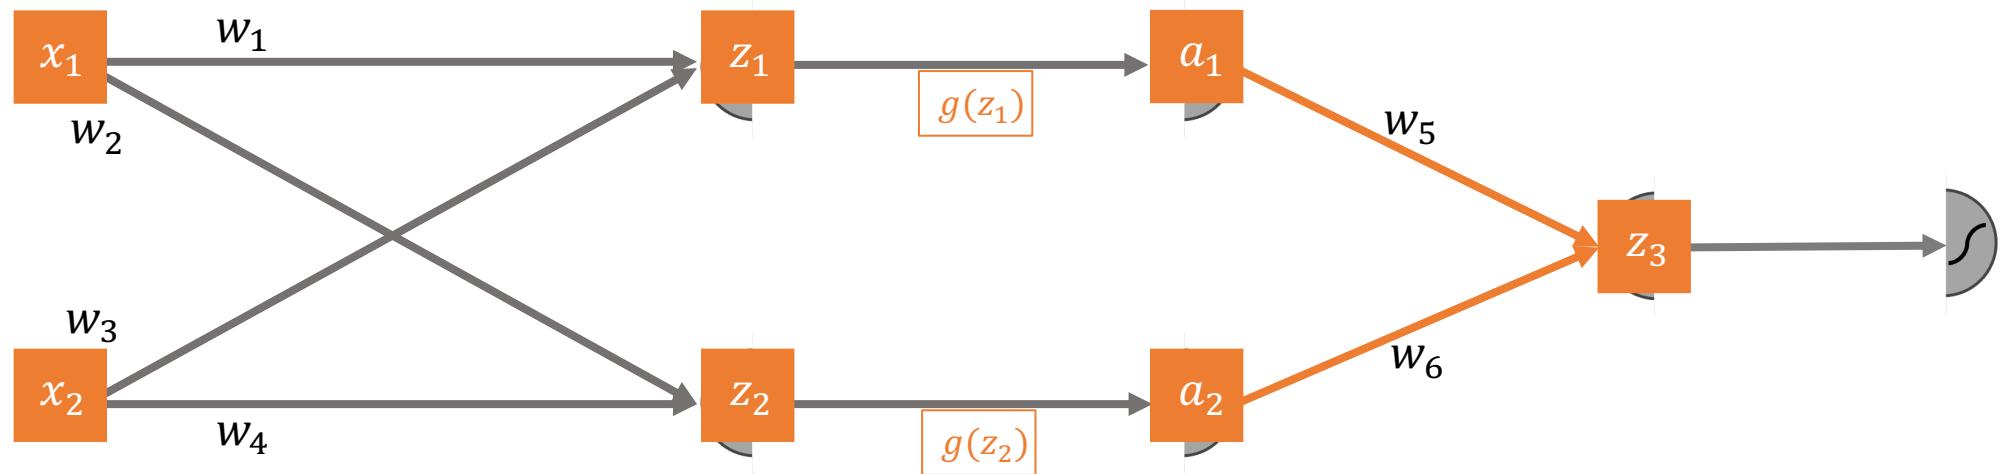
$$a_1 = g(z_1)$$

$$a_2 = g(z_2)$$

Gradient Computation for Multi-layer NN

Activation function: $g(z)$

1. Generate the predicted value \hat{y} for a given data point (x, y) :



$$z_1 = w_1 x_1 + w_3 x_2$$

$$z_2 = w_2 x_1 + w_4 x_2$$

$$a_1 = g(z_1)$$

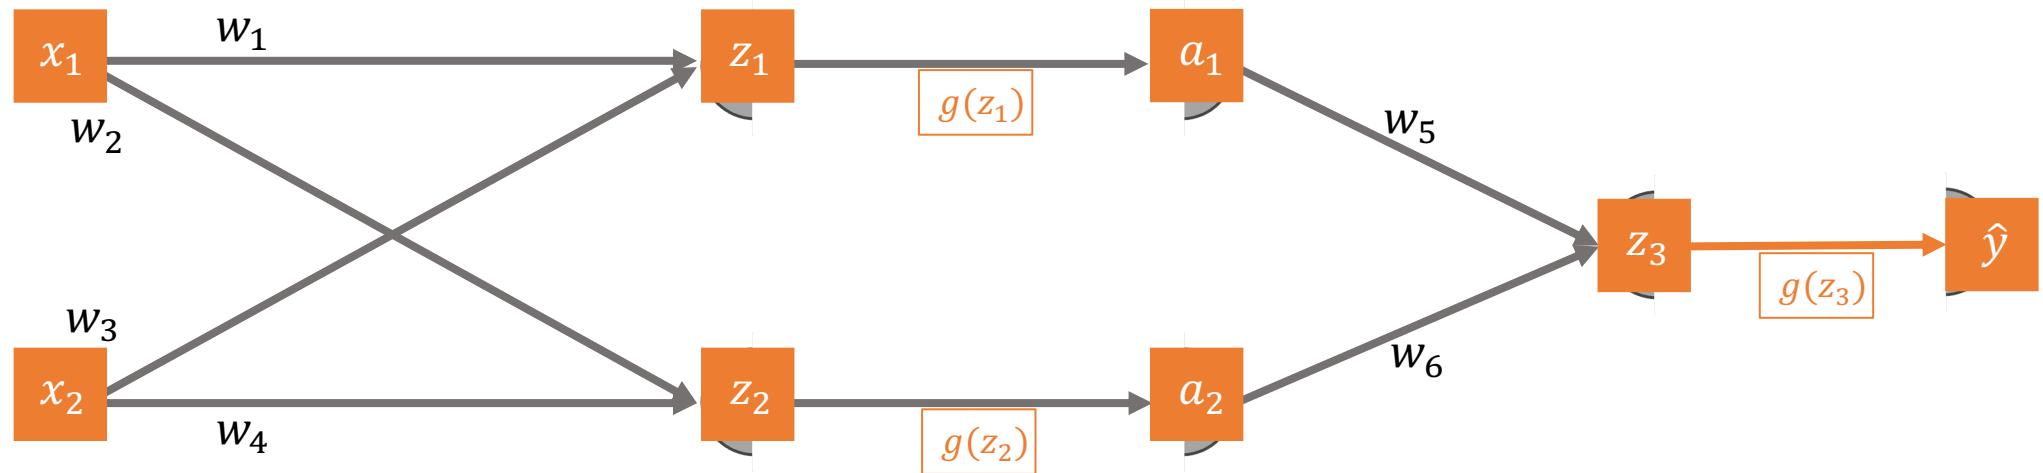
$$a_2 = g(z_2)$$

$$z_3 = w_5 a_1 + w_6 a_2$$

Gradient Computation for Multi-layer NN

Activation function: $g(z)$

1. Generate the predicted value \hat{y} for a given data point (x, y) :



$$z_1 = w_1 x_1 + w_3 x_2$$

$$z_2 = w_2 x_1 + w_4 x_2$$

$$a_1 = g(z_1)$$

$$a_2 = g(z_2)$$

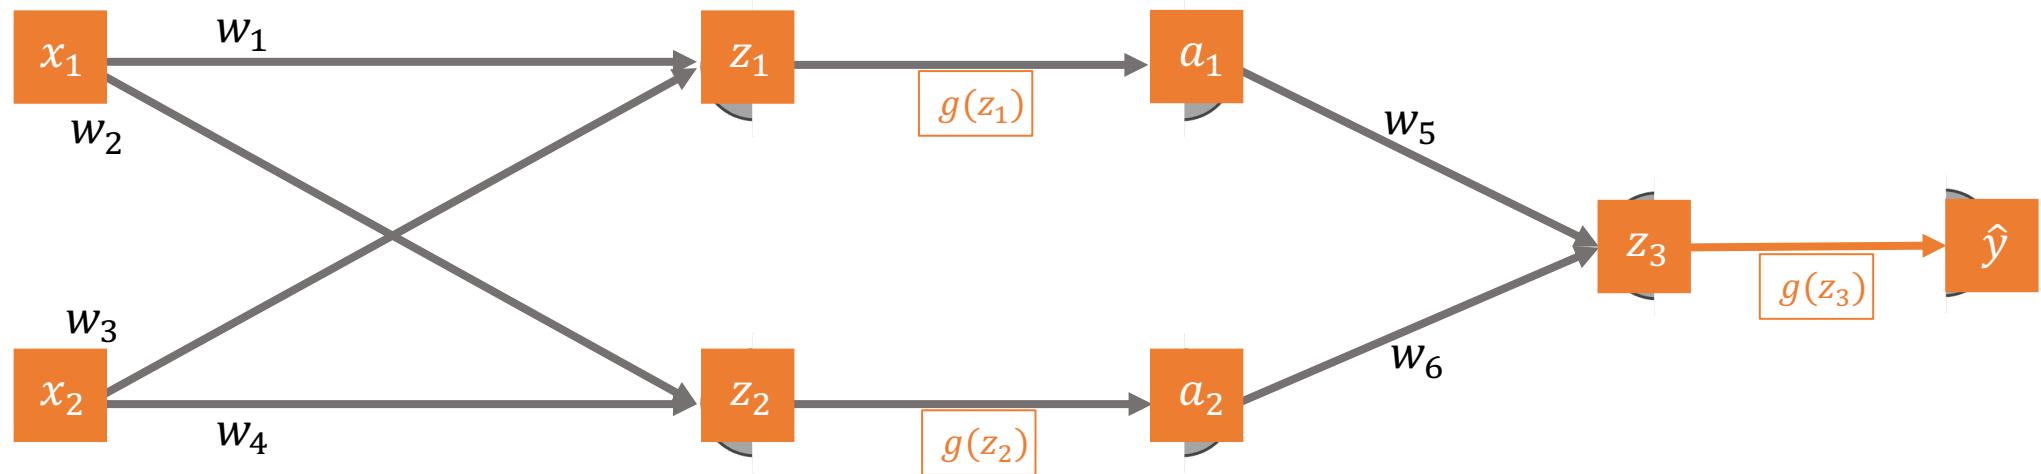
$$z_3 = w_5 a_1 + w_6 a_2$$

$$\hat{y} = g(z_3)$$

Gradient Computation for Multi-layer NN

2. Compute the loss, e.g., MSE

Activation function: $g(z)$
Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



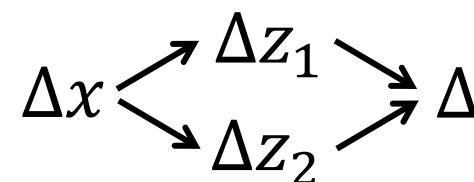
Background: Chain Rule (2)

- The chain rule can also be used to compute the derivative of l :

$$l = h(f(x), g(x))$$

- By introducing the intermediate variables, we can rewrite the l :
Let $z_1 = f(x)$ and $z_2 = g(x)$, then $l = h(z_1, z_2)$

- The chain rule states that the derivative of l with respect to x is given by



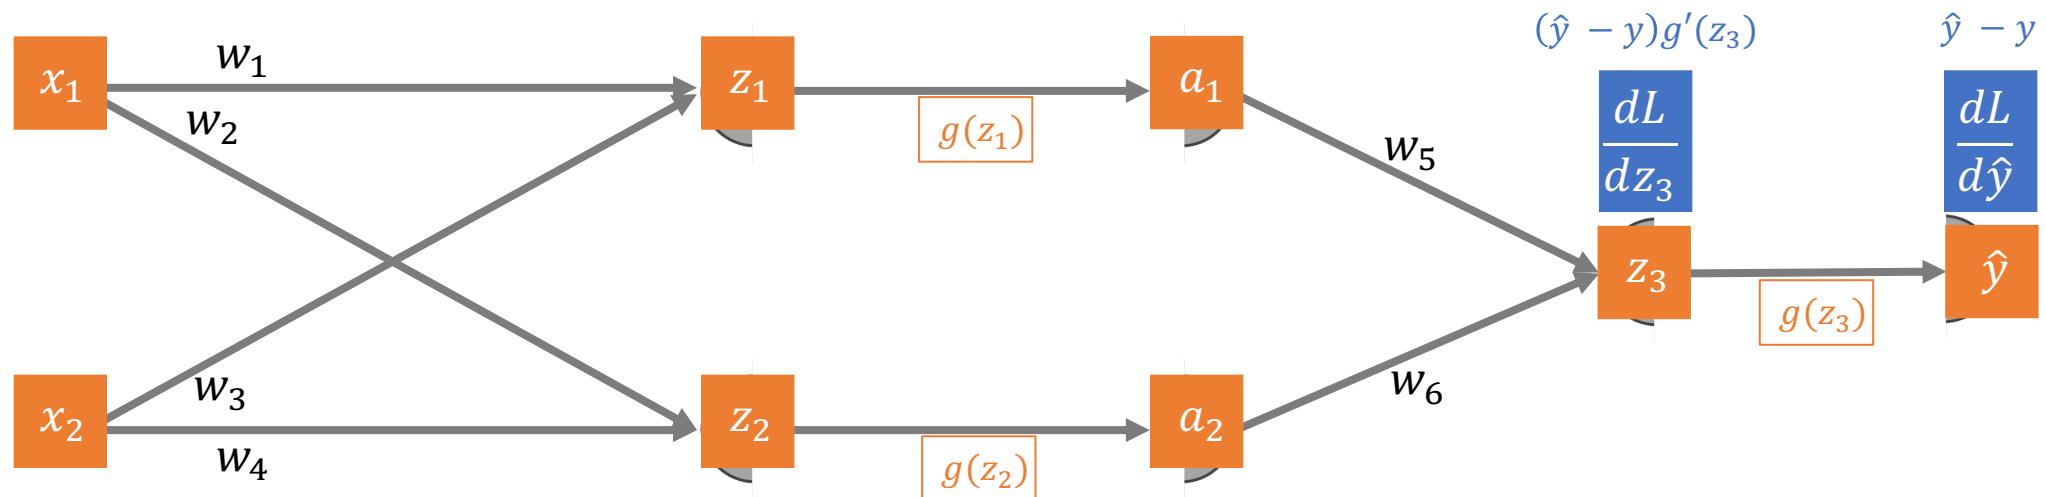
$$\frac{dl}{dx} = \frac{\partial l}{\partial z_1} \frac{dz_1}{dx} + \frac{\partial l}{\partial z_2} \frac{dz_2}{dx}$$

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

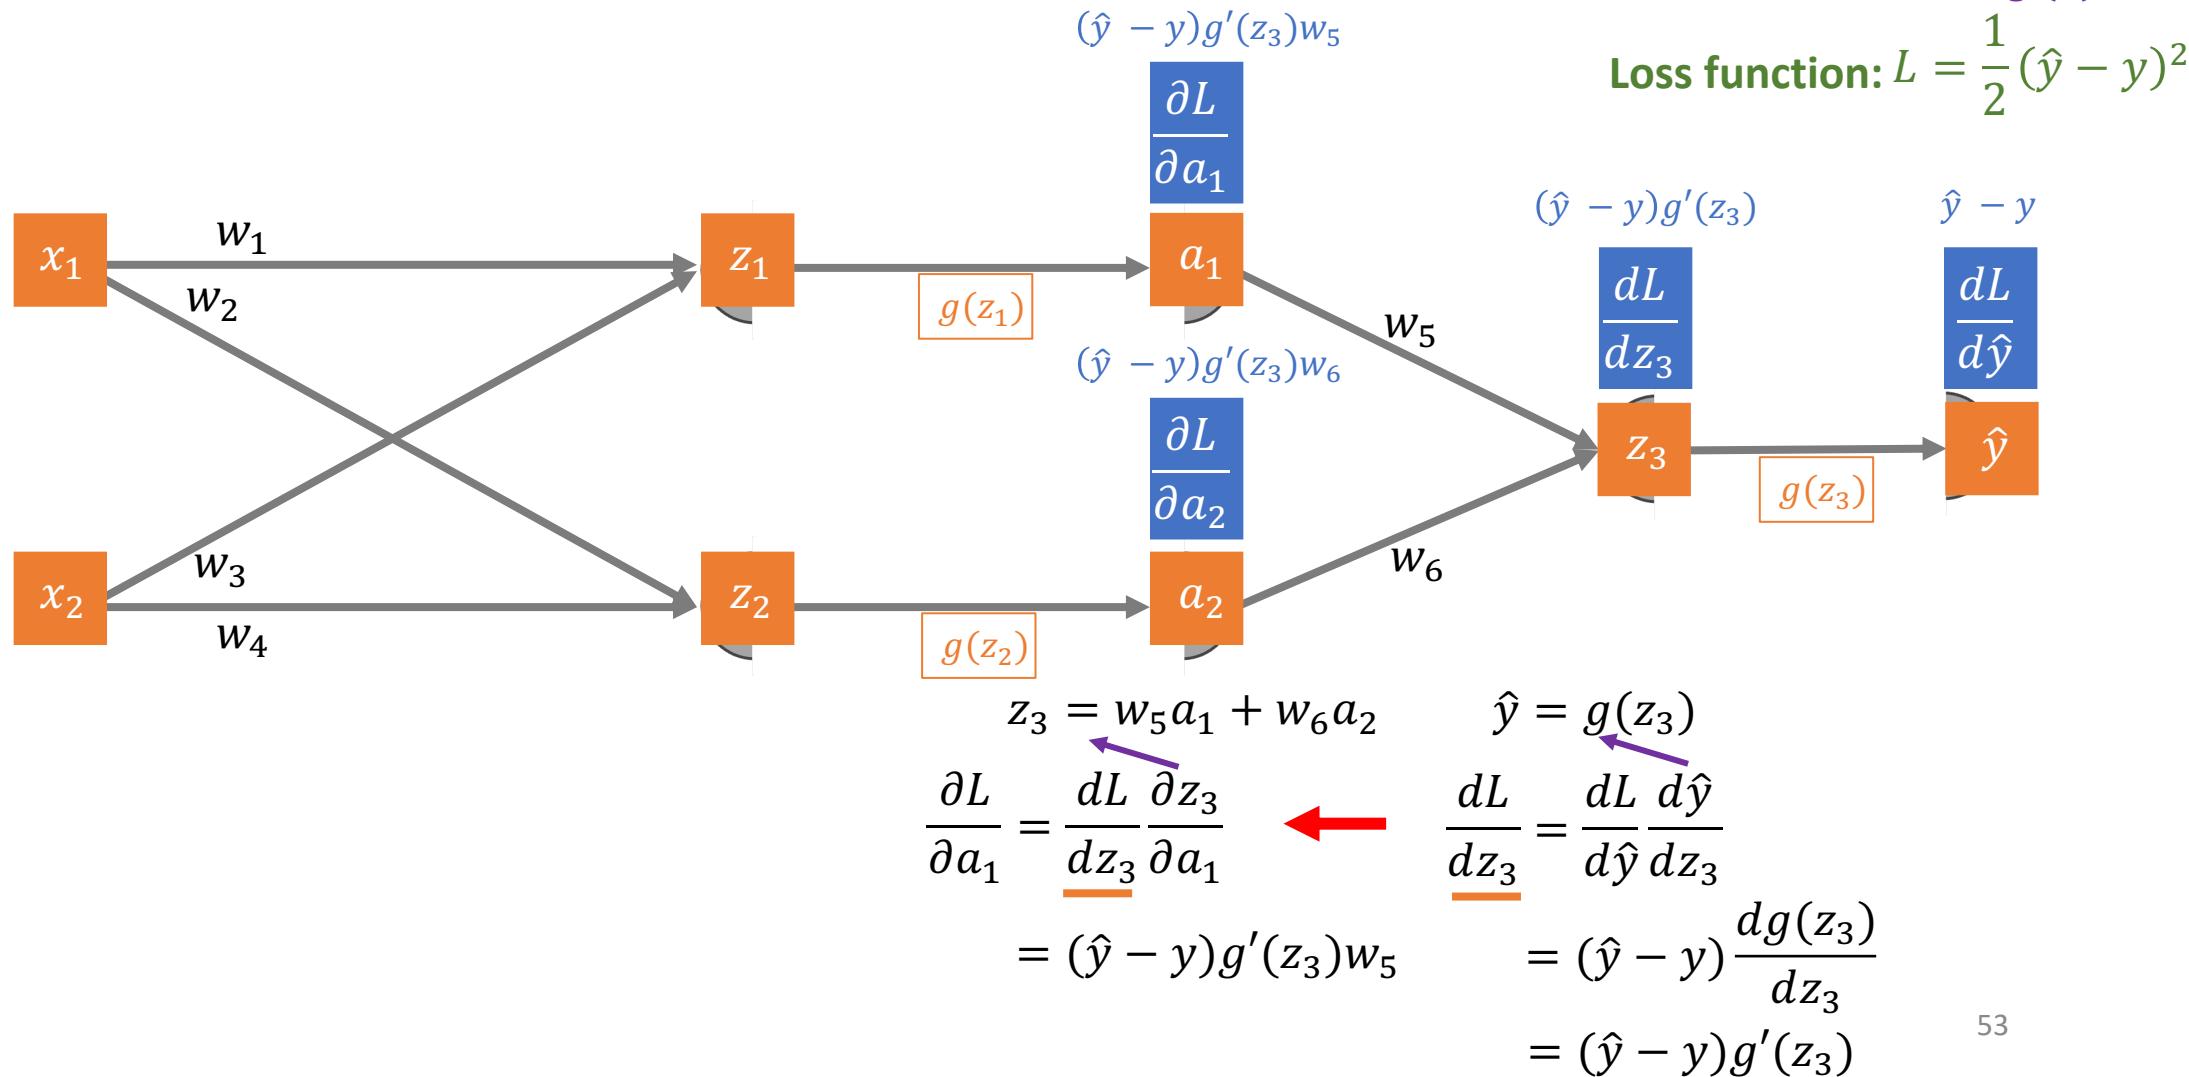
Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



$$\begin{aligned}
 \hat{y} &= g(z_3) \\
 \frac{dL}{dz_3} &= \frac{dL}{d\hat{y}} \frac{d\hat{y}}{dz_3} \\
 &= (\hat{y} - y) \frac{d}{dz}g(z_3) \\
 &= (\hat{y} - y)g'(z_3)
 \end{aligned}$$

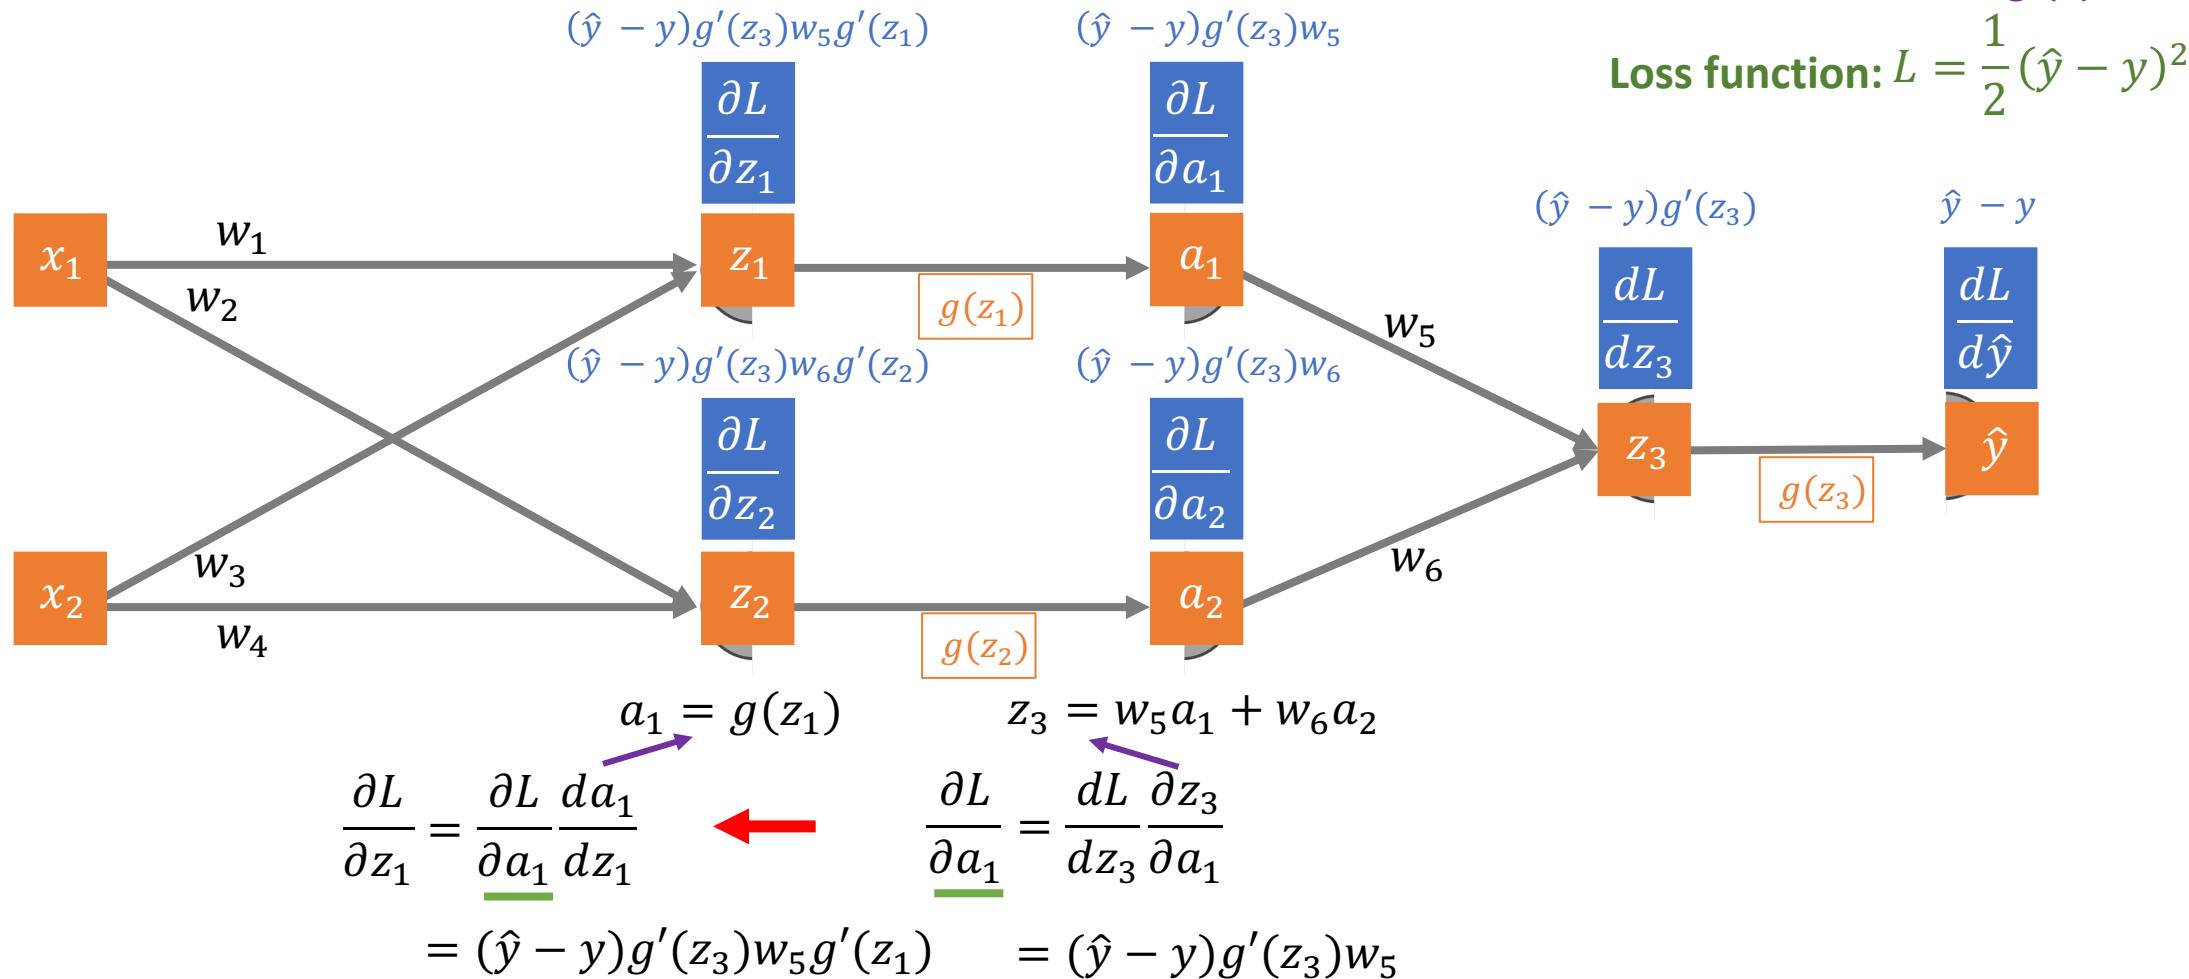
Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation



Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

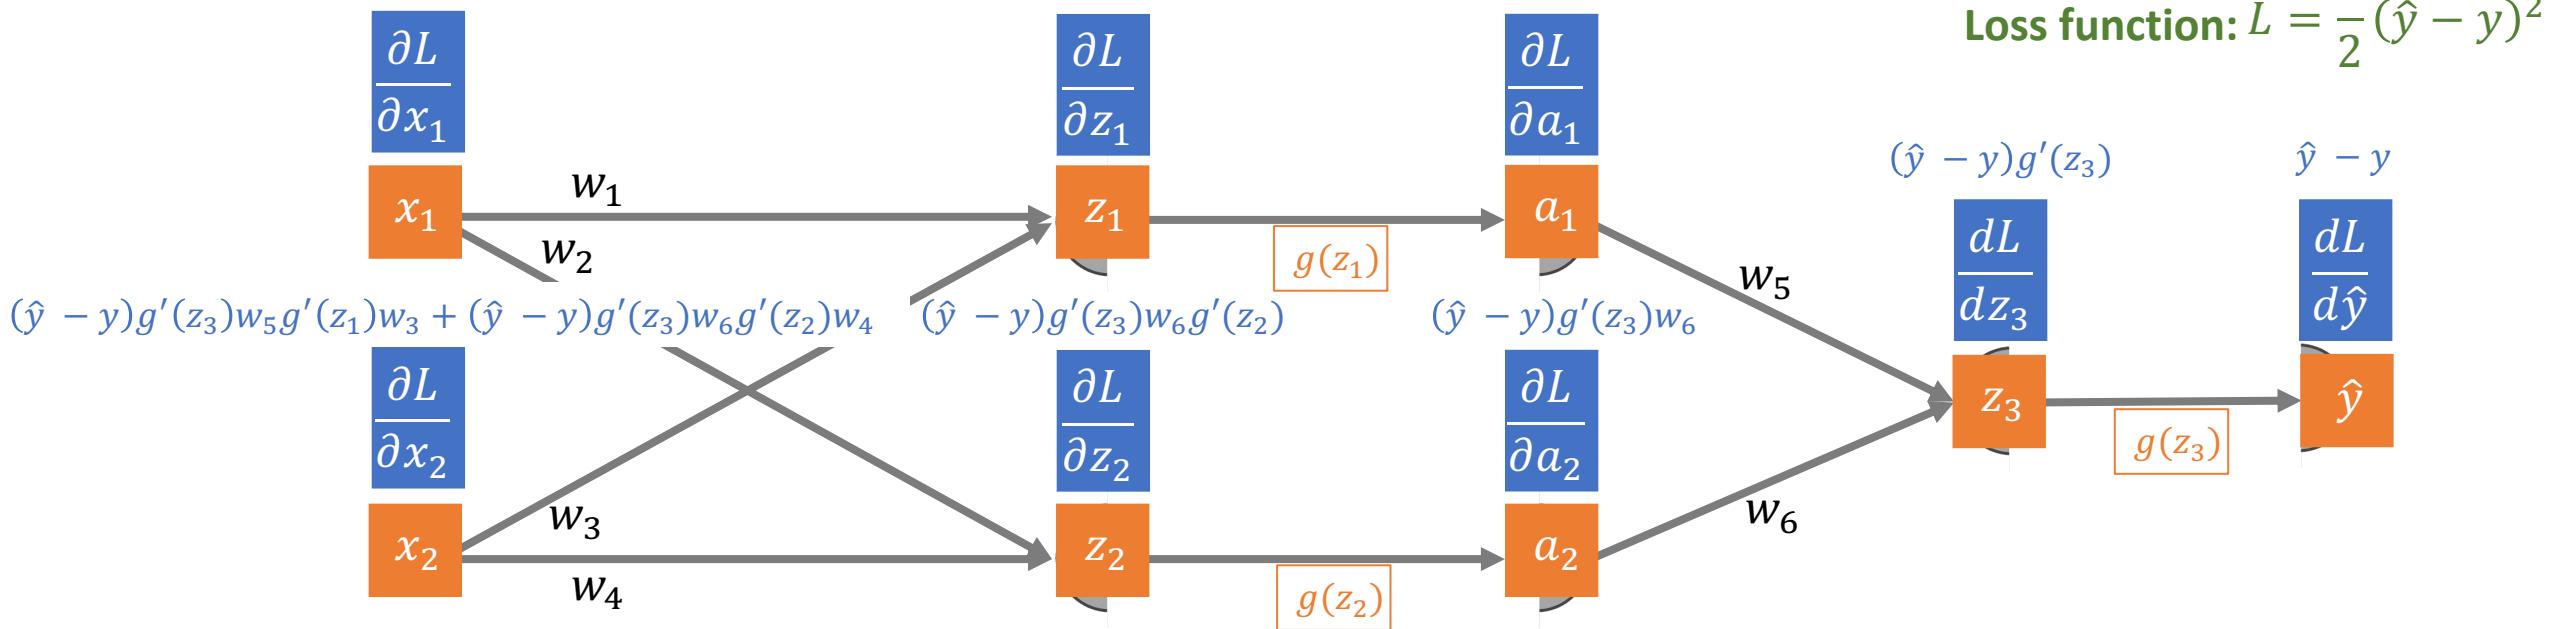


Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

$$(\hat{y} - y)g'(z_3)w_5g'(z_1)w_1 + (\hat{y} - y)g'(z_3)w_6g'(z_2)w_2 \quad (\hat{y} - y)g'(z_3)w_5g'(z_1)$$

$$(\hat{y} - y)g'(z_3)w_5$$



Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$

$$\hat{y} - y$$



$$z_1 = w_1x_1 + w_3x_2 \quad z_2 = w_2x_1 + w_4x_2 \quad a_1 = g(z_1)$$

$$\frac{\partial L}{\partial x_1} = \frac{\partial L}{\partial z_1} \frac{\partial z_1}{\partial x_1} + \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial x_1}$$

$$= (\hat{y} - y)g'(z_3)w_5g'(z_1)w_1 + (\hat{y} - y)g'(z_3)w_6g'(z_2)w_2$$

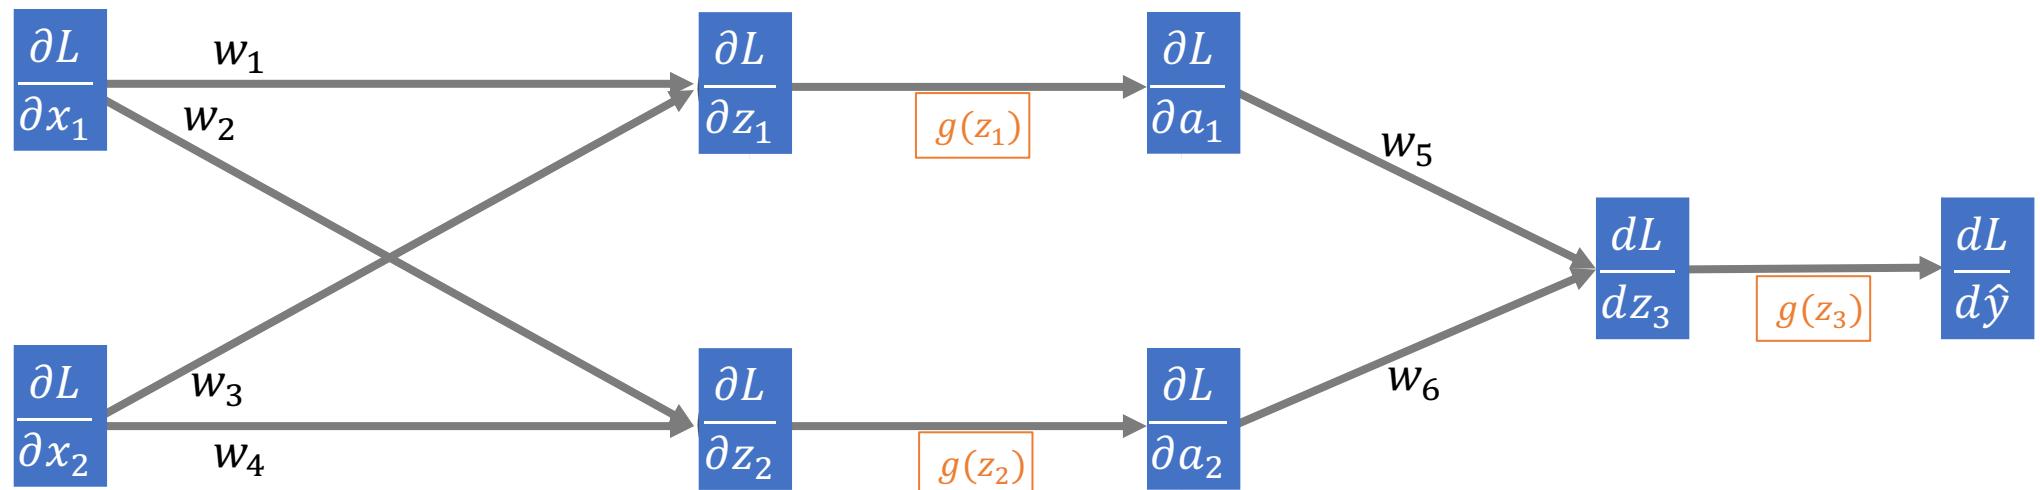
$$= (\hat{y} - y)g'(z_3)w_5g'(z_1)$$

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



A simple way to compute the gradients with respect to values generated in the forward propagation:

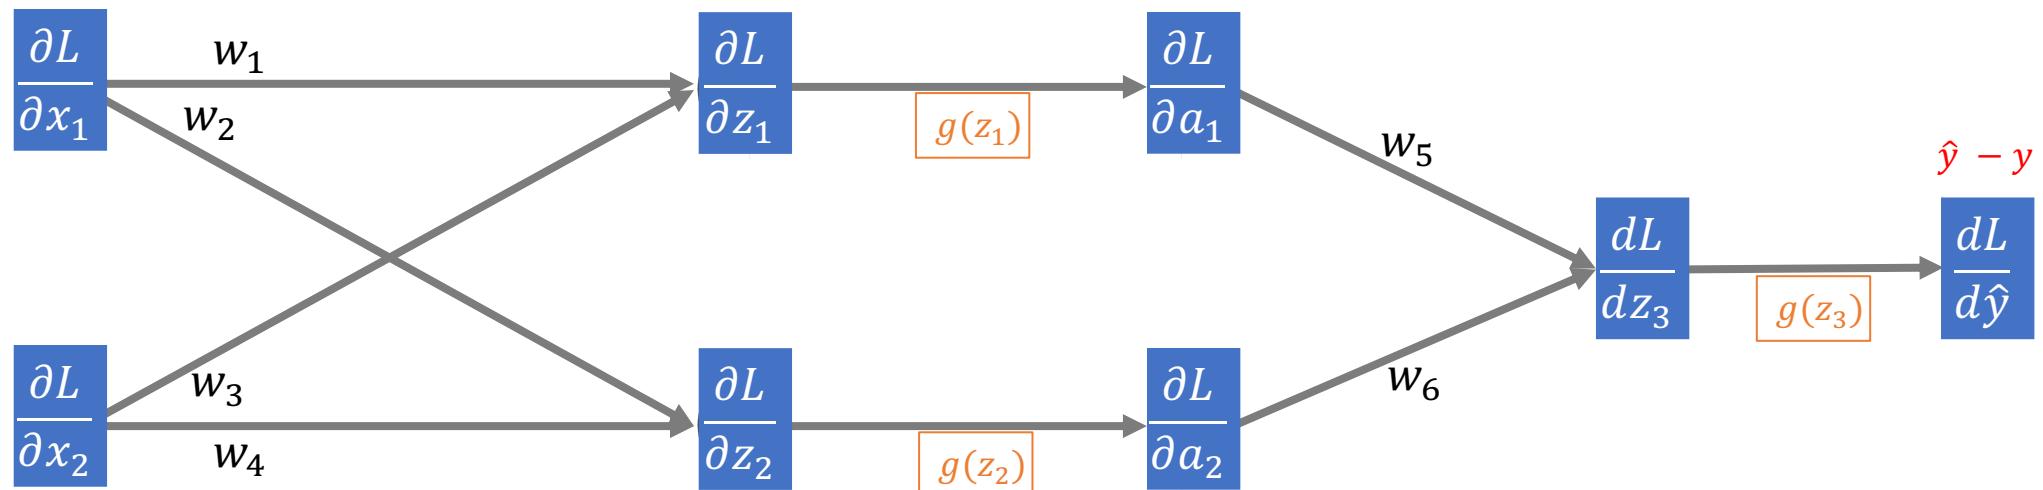
1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. Reverse the graph and compute gradients

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



A simple way to compute the gradients with respect to values generated in the forward propagation:

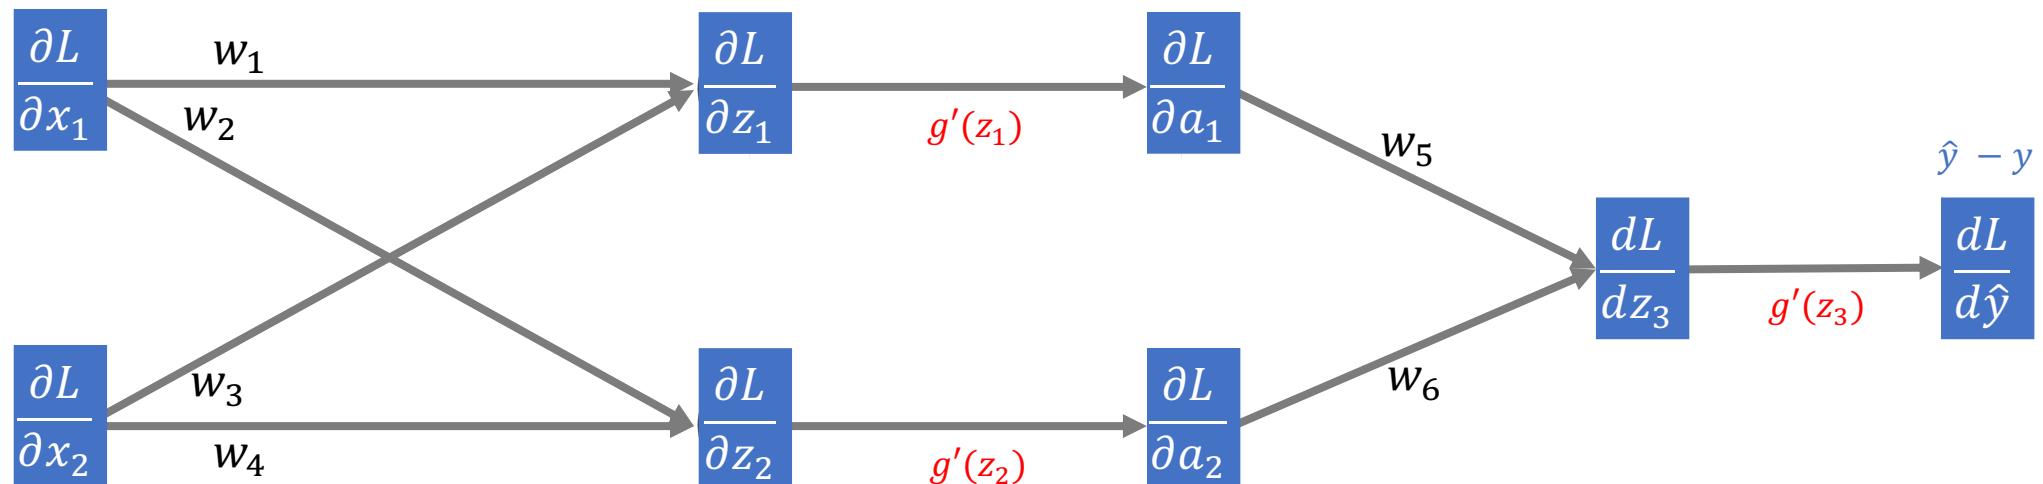
1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. Reverse the graph and compute gradients

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



A simple way to compute the gradients with respect to values generated in the forward propagation:

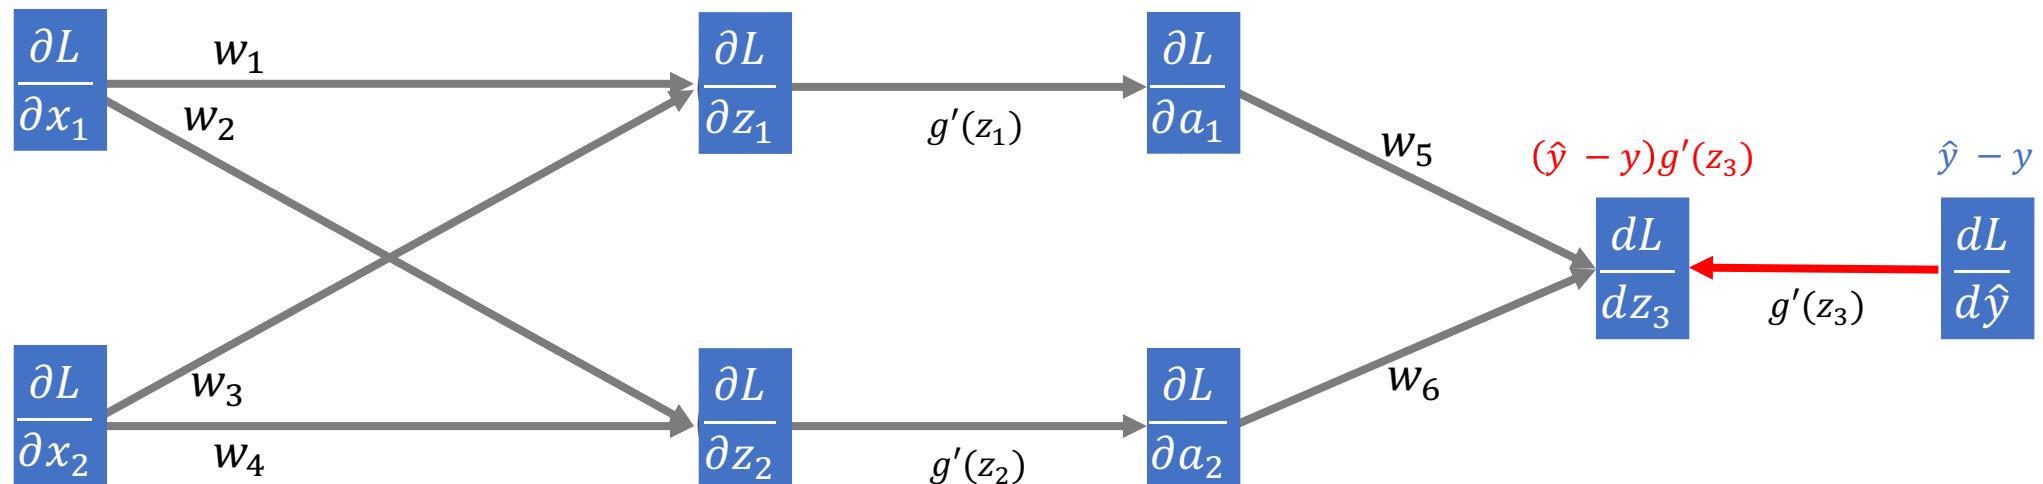
1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. Reverse the graph and compute gradients

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



A simple way to compute the gradients with respect to values generated in the forward propagation:

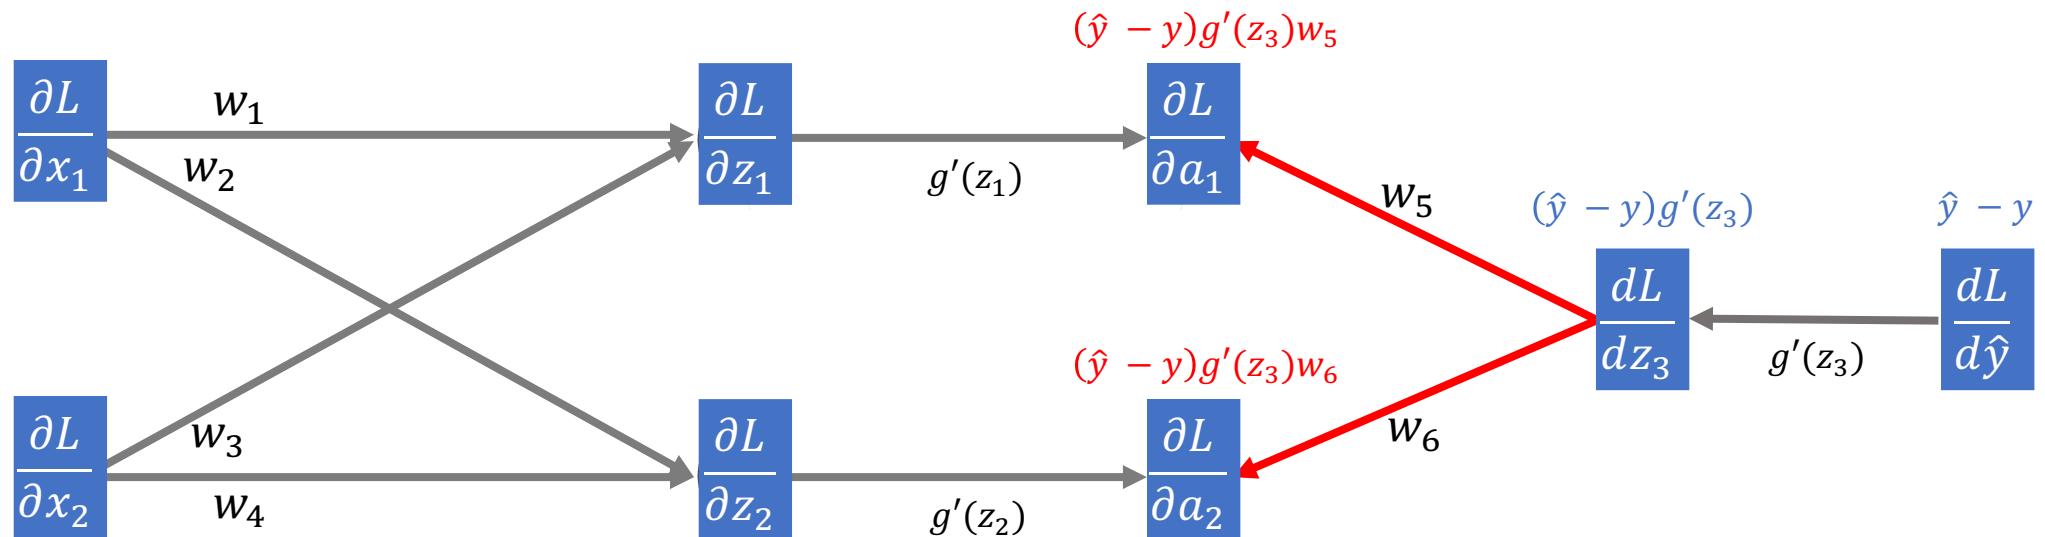
1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. **Reverse the graph and compute gradients**

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



A simple way to compute the gradients with respect to values generated in the forward propagation:

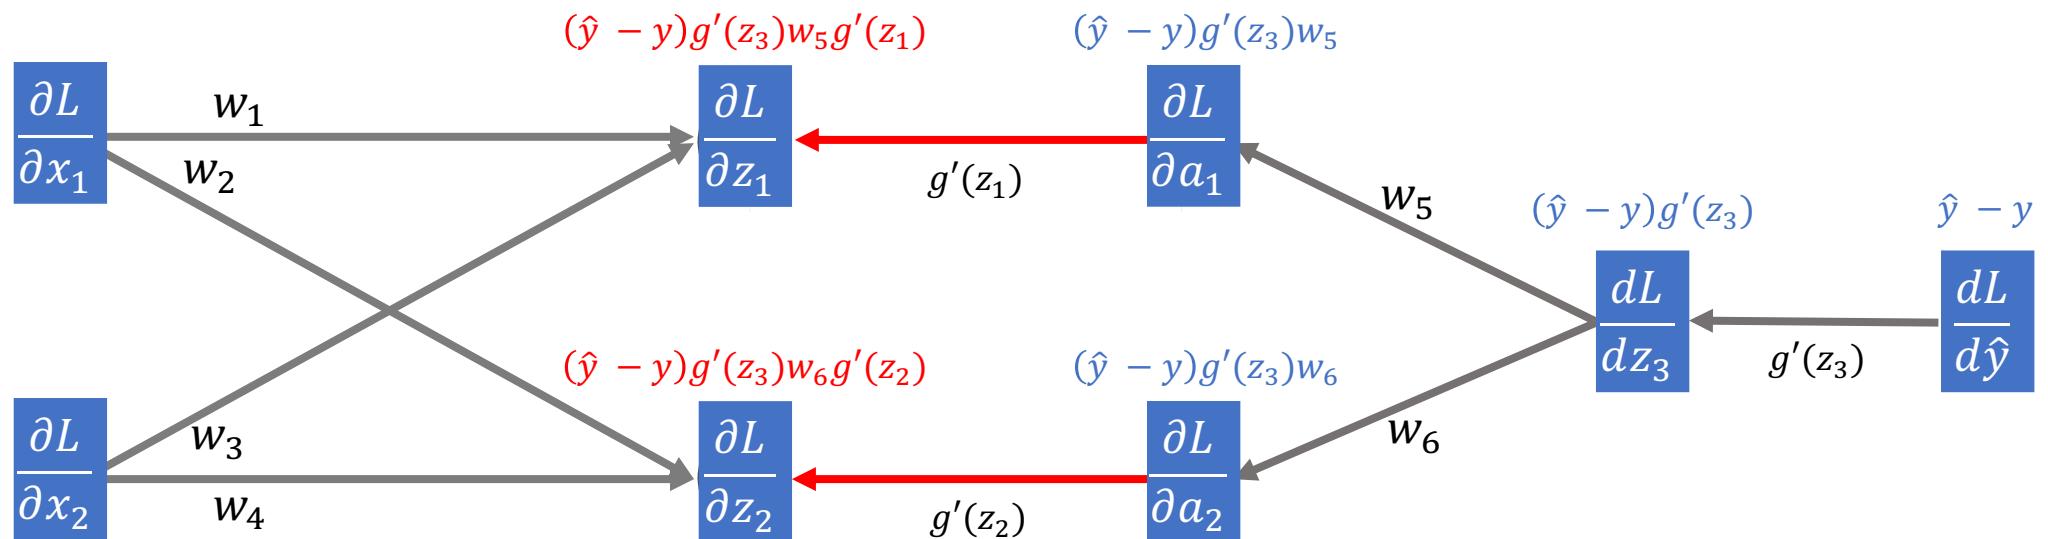
1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. Reverse the graph and compute gradients

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

Activation function: $g(z)$

Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$



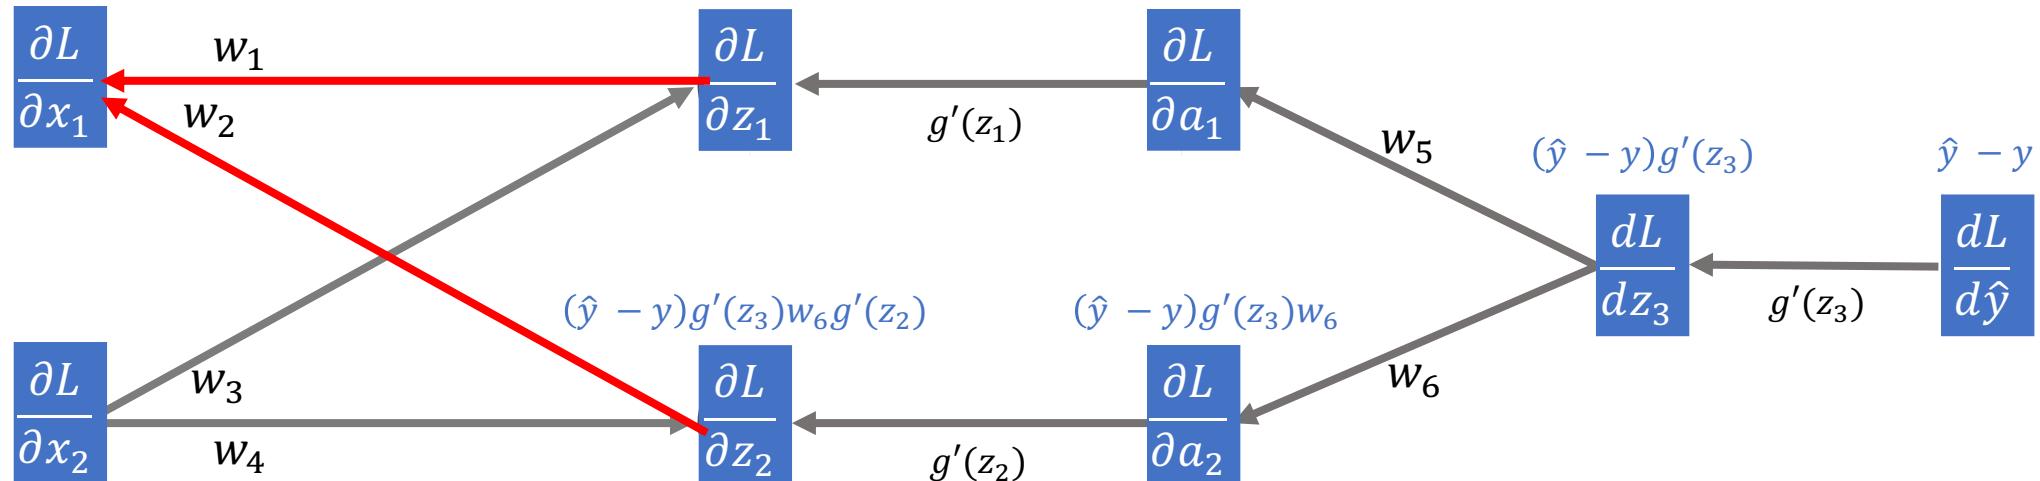
A simple way to compute the gradients with respect to values generated in the forward propagation:

1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. Reverse the graph and compute gradients

Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

$$(\hat{y} - y)g'(z_3)w_5g'(z_1)w_1 \\ + \\ (\hat{y} - y)g'(z_3)w_6g'(z_2)w_2$$



A simple way to compute the gradients with respect to values generated in the forward propagation:

1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. **Reverse the graph and compute gradients**

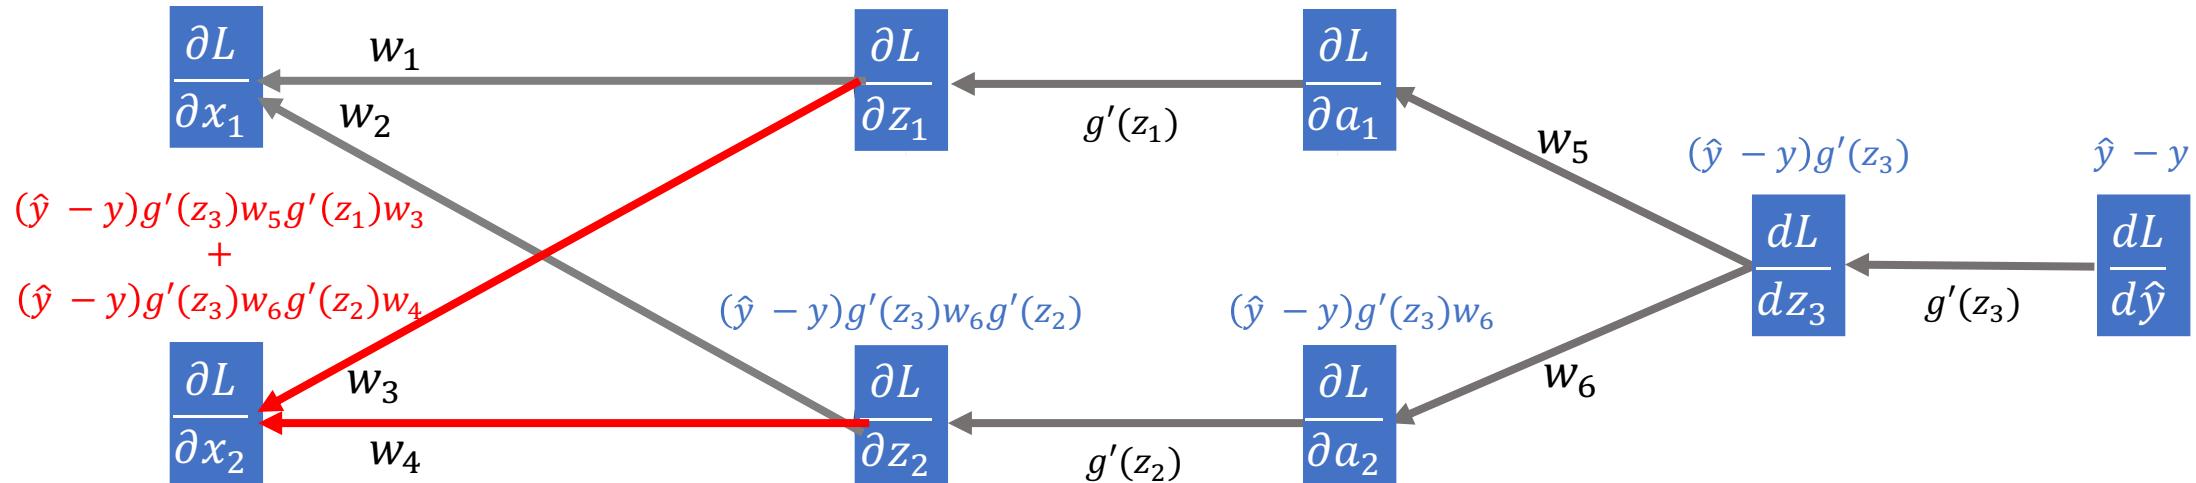
Gradient Computation for Multi-layer NN

Gradients with respect to **values** generated in the forward propagation

$$(\hat{y} - y)g'(z_3)w_5g'(z_1)w_1$$

+

$$(\hat{y} - y)g'(z_3)w_6g'(z_2)w_2$$



Activation function: $g(z)$

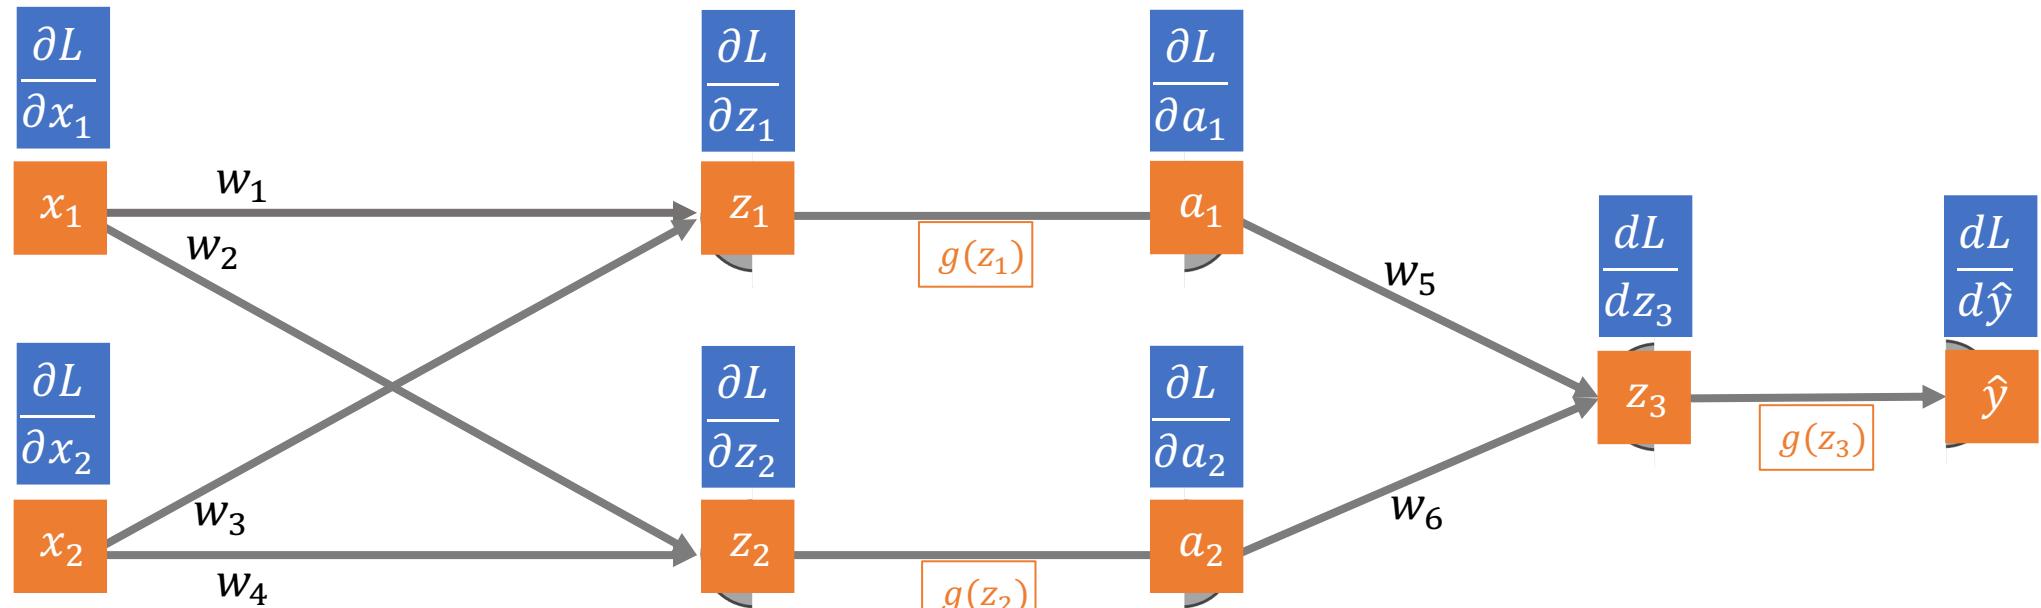
Loss function: $L = \frac{1}{2}(\hat{y} - y)^2$

A simple way to compute the gradients with respect to values generated in the forward propagation:

1. Compute gradient of loss with respect to \hat{y}
2. Convert $g(z)$ to $g'(z)$ and treat $g'(z)$ as the weight associated with the corresponding edge
3. **Reverse the graph and compute gradients**

Gradient Computation for Multi-layer NN

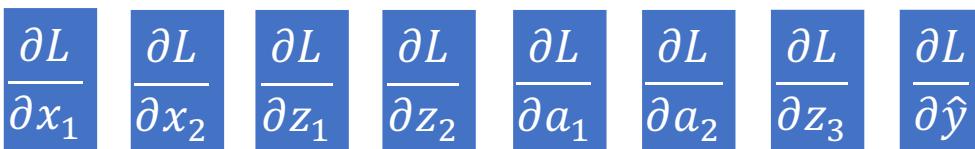
Gradients with respect to **values** generated in the forward propagation



After one forward pass, we can obtain

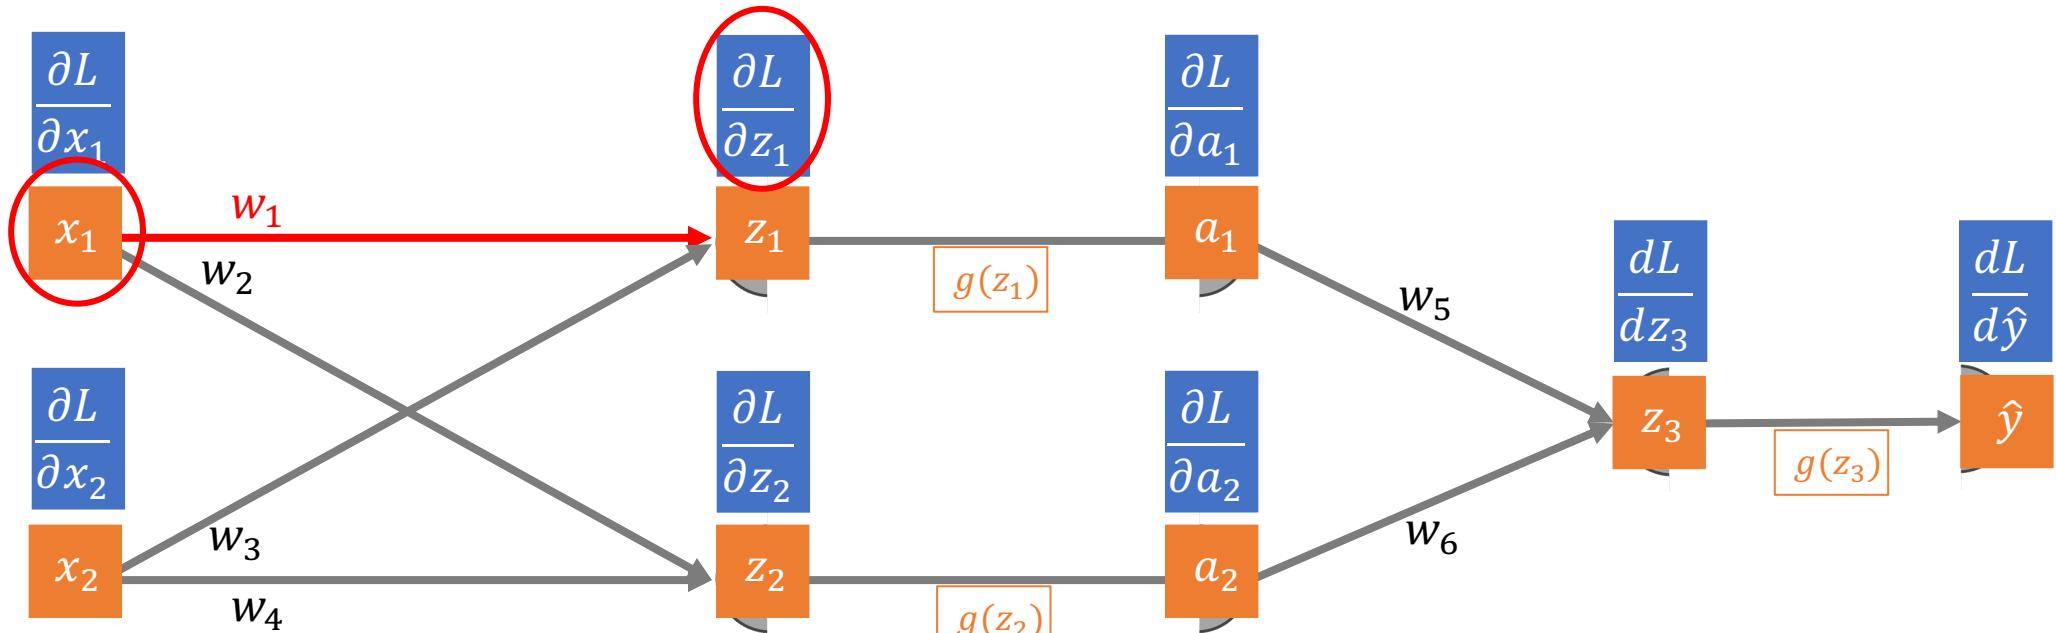


After one backward pass, we can obtain



Gradient Computation for Multi-layer NN

Gradients with respect to **weights**



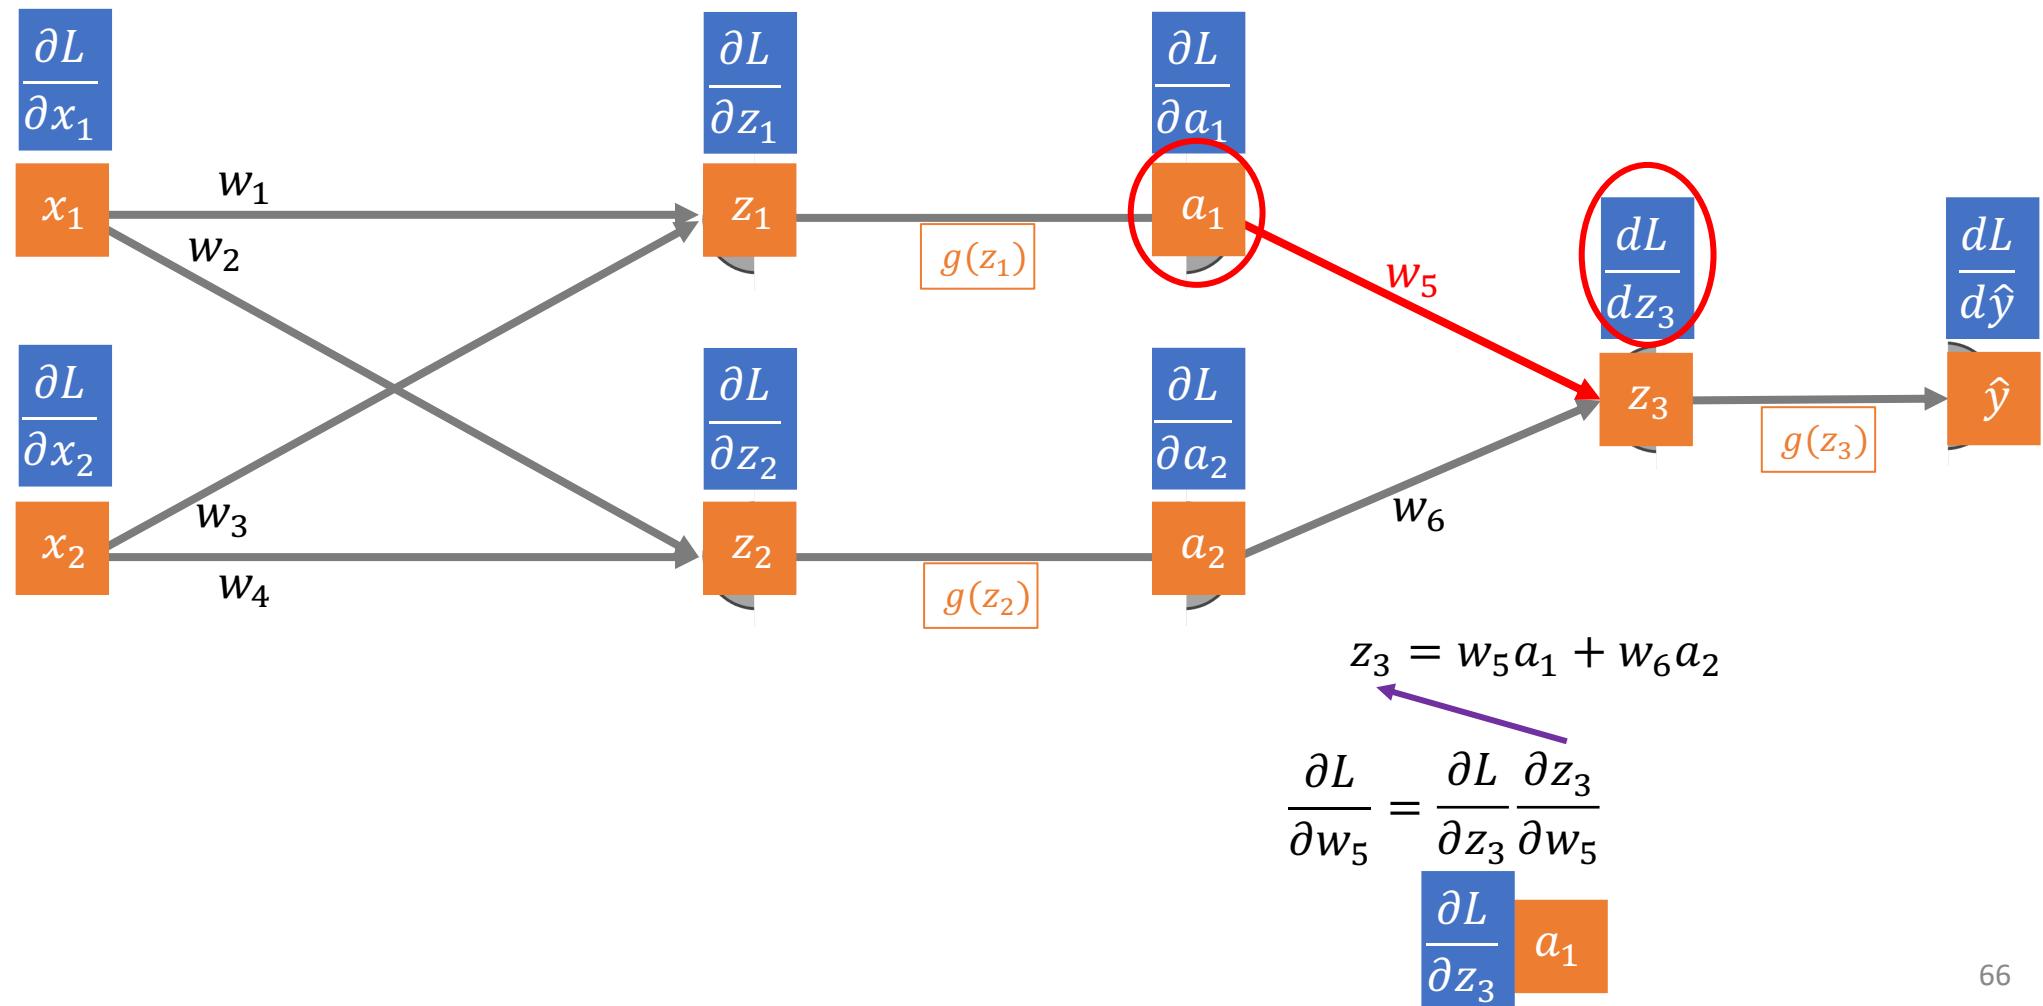
$$z_1 = w_1 x_1 + w_3 x_2$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial z_1} \frac{\partial z_1}{\partial w_1}$$

$$\frac{\partial L}{\partial z_1} x_1$$

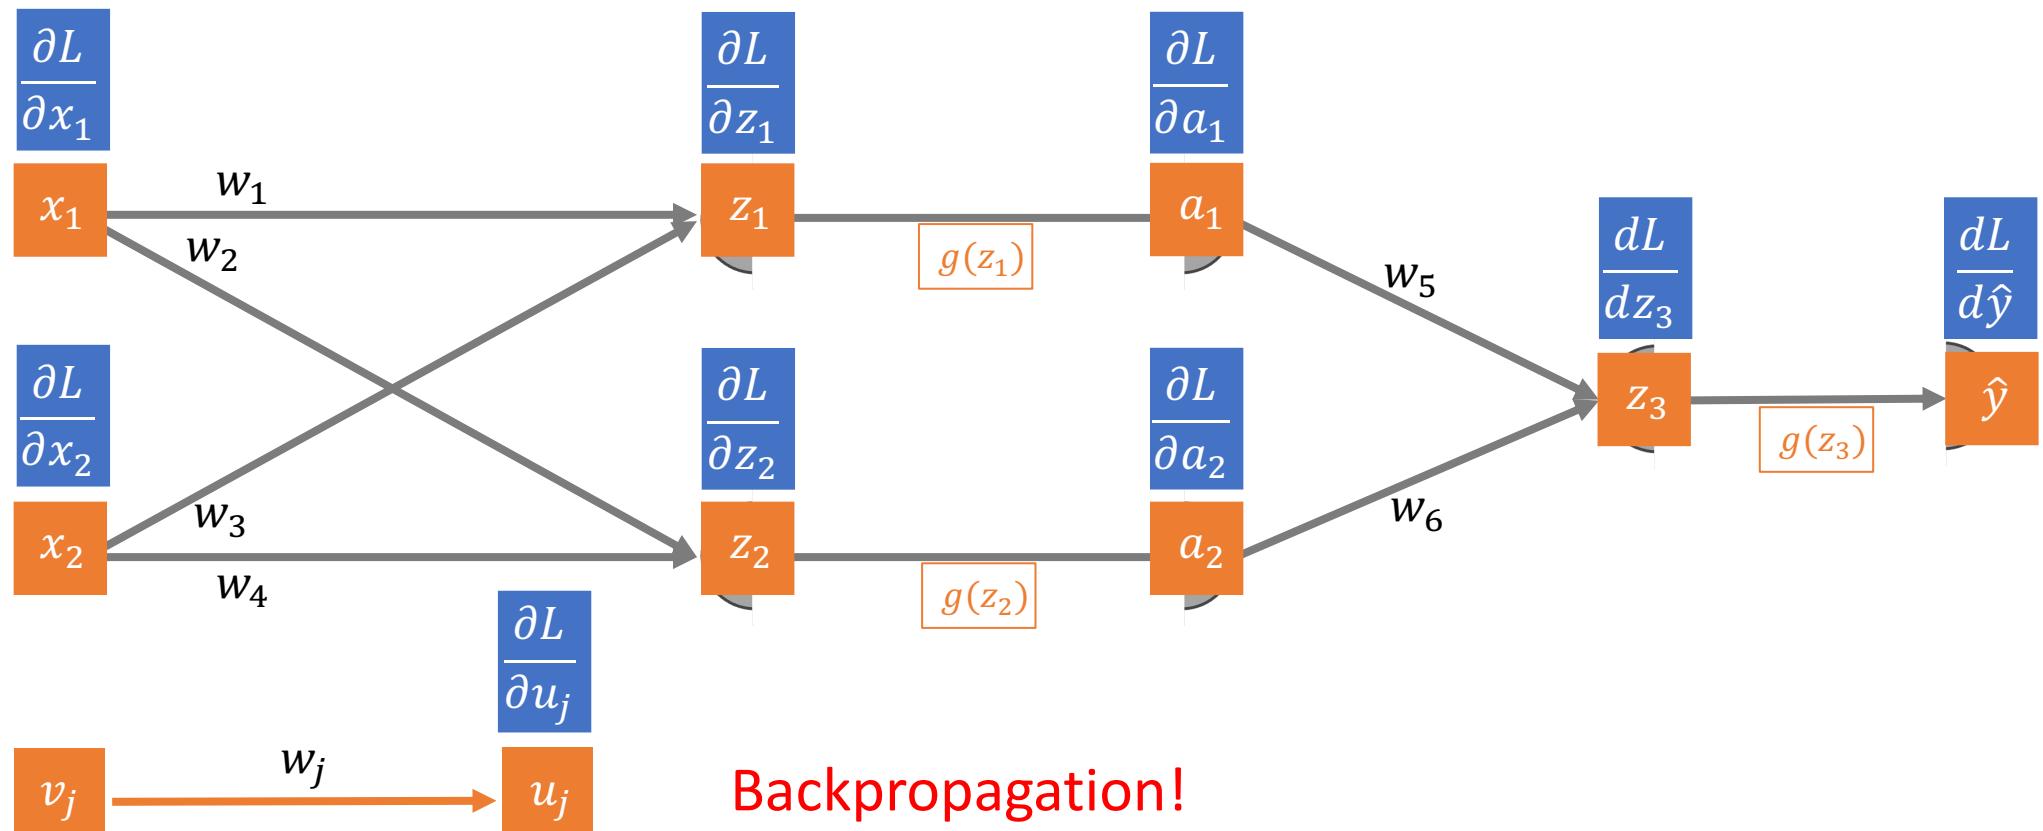
Gradient Computation for Multi-layer NN

Gradients with respect to **weights**



Gradient Computation for Multi-layer NN

Gradients with respect to **weights**



For weight w_j , $\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial u_j} v_j$

$\frac{\partial L}{\partial u_j}$ can be computed in one backward pass.

v_j can be computed in one forward pass.

Summary

- Neural Networks
 - Linear regression model: A neuron with linear activation function
 - Logistic regression model: A neuron with sigmoid activation function
 - Neurons can be used to transform features.
- Tasks for Neural Network
 - Hidden layers are used to learn feature transformation.
 - Output layer should be set according to the target.
- Backpropagation
 - In order to update the weights, we need to compute the gradients.
 - Backpropagation can be used to efficiently compute all the gradients.

Coming Up Next Week

- **More on Neural Networks**
 - Convolutional Neural Networks
 - Recurrent Neural Networks
 - Issues with Deep Learning (If time permit)

To Do

- **Lecture Training 9**
 - +550 EXP
 - +100 Early bird bonus

